Model Evaluation

7B

CS 3244
Machine Learning

NUS National University of Singapore | Computing

[Instructor] Brian Lim

IoT Sensors

Health Behavior Change

Data Analytics

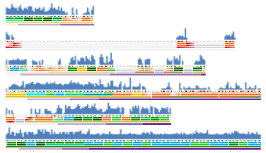

NUS Ubicomp Lab

Apps and Analytics for Smart Cities and Healthcare

http://ubiquitous.comp.nus.edu.sg

Explainable Artificial Intelligence

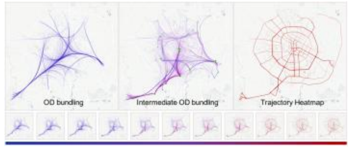

Interactive Data Visualization

brianlim@comp.nus.edu.sg

Academic Experience

- Asst. Prof. in Computer Science
- Ph.D. in HCI, Carnegie Mellon University
- B.S. in Engineering Physics, Cornell University

Research Interests

- HCI: understand people with tech, help people with tech
- Explainable Artificial Intelligence
- Ubiquitous Computing
- Data analysis and visualization
- Smart Health and Smart Cities

Recap from Past Weeks

Three learning problems

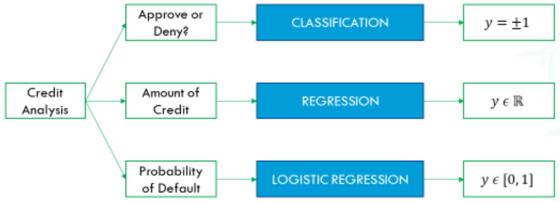

Linear models are the fundamental models.

The linear model is the first model to try.

NUS CS3244: Machine Learning 14

W04

During COVID-19 Pandemic

Are you sick?

$$f \longrightarrow \begin{cases} +1 & \text{sick} \\ -1 & \text{well} \end{cases}$$

False negative highly costly!
People and the economy dies.

False positive requires the inconvenience of quarantine.

h_θ \ f	+1	−1
+1	0 True Positive	1 False Positive
−1	10,000 False Negative	0 True Negative

NUS CS3244: Machine Learning 35

W03

Analyzing the estimate

On a test <u>point</u> (x, y), the cost $J(h_\theta(x), y)$ is:

Squared error: $(h_\theta(x) - y)^2$

Binary error: $[[h_\theta(x) \neq y]]$

$$\mathbb{E}[J(h_\theta(x), y)] = L_{test}(h_\theta)$$

$$Var[J(h_\theta(x), y)] = \sigma^2$$

NUS CS3244: Machine Learning

W06

Distance Metric

Often we'll use Euclidean distance as our first try:

$$L2(x^{(*)}, x^{(i)}) = \sqrt{\left(x_1^{(*)} - x_1^{(i)}\right)^2 + \left(x_2^{(*)} - x_2^{(i)}\right)^2 + \cdots + \left(x_n^{(*)} - x_n^{(i)}\right)^2}$$

but any *distance metric* (*non-negativity, symmetric, obeys the triangle inequality*) is possible.

Defining your distance metric to fit your data's semantic meaning can be important.

NUS CS3244: Machine Learning 65

W02

Have your cake and eat it too: K fold cross validation

10 FOLD C.V.

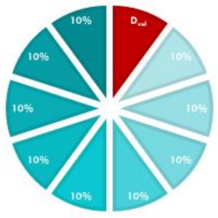

LOOCV can be very expensive for large datasets. Why?

Instead, use K fold cross validation: i.e., K training sessions on $\frac{m}{K}$ points each.

Recommend: 10-fold CV

NUS CS3244: Machine Learning 51

W06

Week 07b: Learning Outcomes

- Describe various evaluation metrics of model performance
- Understand that model performance depends on prediction task and data
  - Describe several challenges in evaluating model performances
  - Choose appropriate **evaluation metric** for different prediction tasks

Week 07b: Lecture Outline

1. Recap: Supervised learning Classification vs. Regression
2. Classification Metrics
3. Regression Metrics [W08a]
4. Unsupervised learning metrics [W11]

Classification vs. Regression

NUS National University of Singapore | Department of Computer Science
School of Computing

<table>
<tr>
<td>

Classification

$y \in \{0,1\}$ binary

$y \in \{y_A, y_B, \dots\}$ multi-class

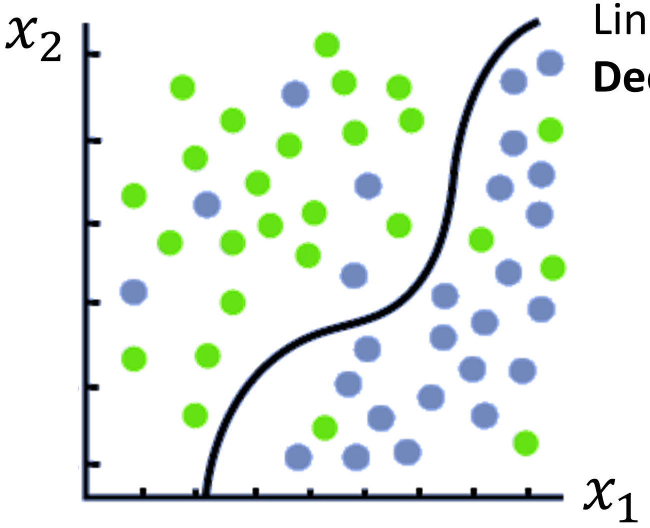

Line is a **Decision Boundary**

</td>
<td>

Note different axes! The are not the same type

</td>
<td>

Regression

$y \in \mathbb{R}$ any real number

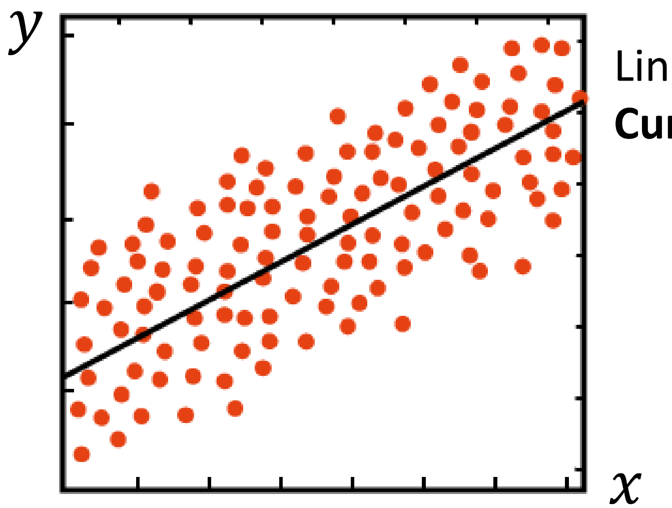

Line is a **Curve Fit**

</td>
</tr>
</table>

Image credit: https://www.javatpoint.com/regression-vs-classification-in-machine-learning

Classification

$y \in \{0,1\}$ binary

$y \in \{y_A, y_B, \dots\}$ multi-class

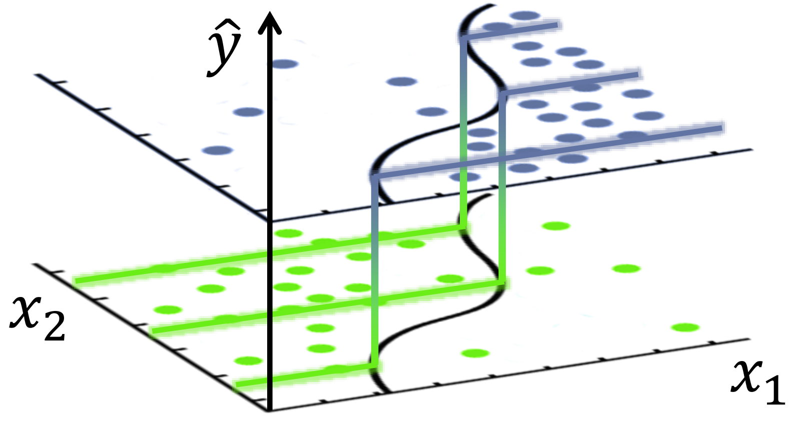

$$y = M(\boldsymbol{x}), \quad \boldsymbol{x} = \vec{x} = (x_1, x_2)^\top$$

Regression

$y \in \mathbb{R}$ any real number

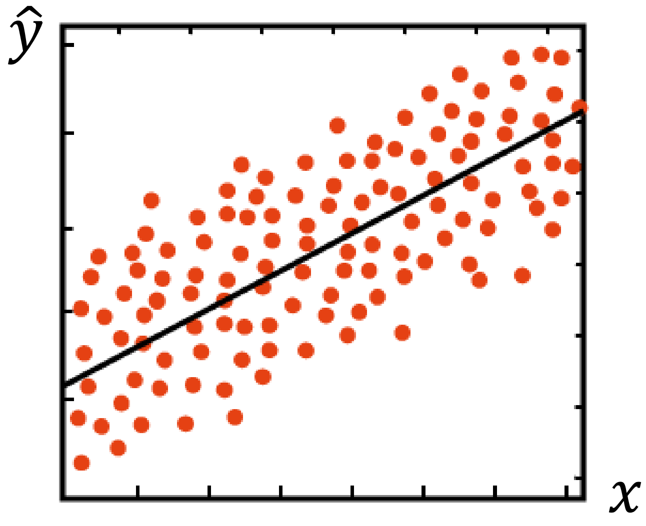

$$y = M(x), \quad x = x_1$$

Image credit:
https://www.javatpoint.com/regression-vs-classification-in-machine-learning

Classification

$y \in \{0,1\}$ binary

$y \in \{y_A, y_B, \dots\}$ multi-class

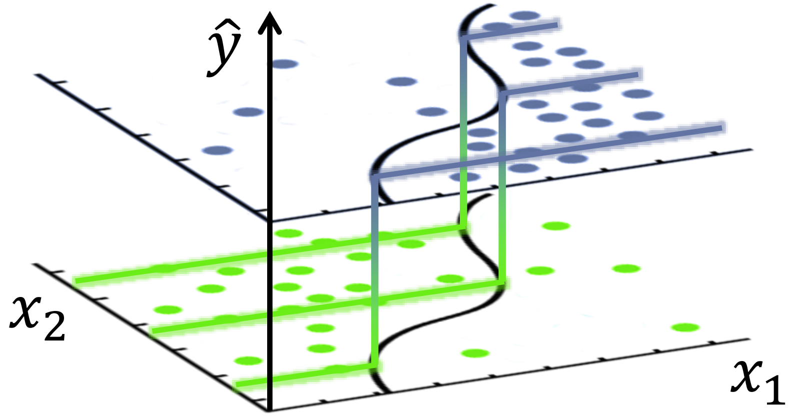

$$y = M(\boldsymbol{x}), \quad \boldsymbol{x} = \vec{x} = (x_1, x_2)^\top$$

Regression

$y \in \mathbb{R}$ any real number

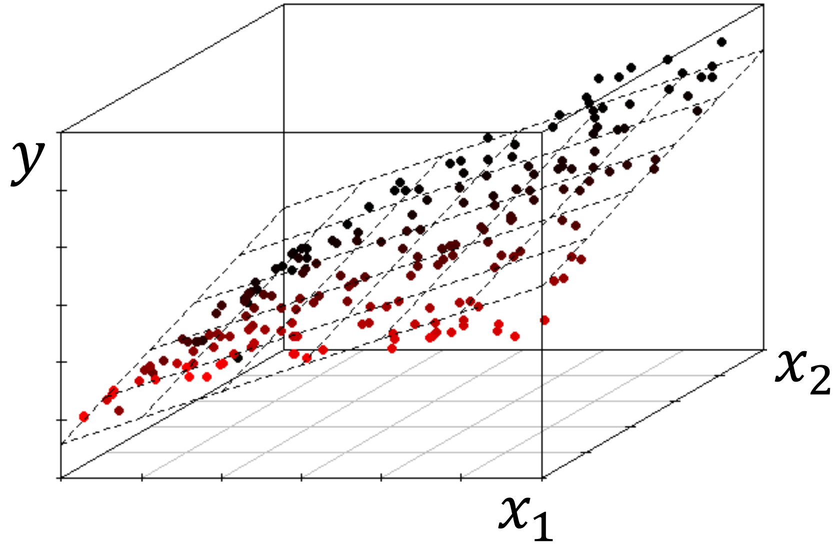

$$y = M(x), \quad x = x_1$$

$$y = M(\boldsymbol{x}), \quad \boldsymbol{x} = (x_1, \dots, x_n)^\top$$

Image credit:
https://www.javatpoint.com/regression-vs-classification-in-machine-learning,
https://stackoverflow.com/q/26431800

Classification Evaluation Metrics

NUS National University of Singapore | Department of Computer Science School of Computing

Week 07: Lecture Outline

1. Recap: Classification vs. Regression
2. Classification Metrics
    1. Accuracy
    2. Confusion Matrix, TP, TN, FP, FN
    3. Precision, Recall, F_1
    4. ROC, AUC
    5. Micro- and Macro-Averaging
    6. PR-AUC (Average Precision)

    Performance Metrics
3. Regression Metrics [W08a]

Correctness

Classification is correct when prediction $\hat{y}$ is the same as actual label y, i.e.,

$$\text{Correct} = [\hat{y} = y]$$

where

- $\hat{y} = M(\boldsymbol{x})$ is the predicted value from model M instance $\boldsymbol{x}$
- y is the ground truth value
- $[P] = \begin{cases} 1 & \text{if } P \text{ is true} \\ 0 & \text{otherwise} \end{cases}$ is the Iverson bracket notation for if/else

Accuracy

"Average correctness" across test dataset with m instances:

$$A = \frac{1}{m} \sum_{j=1}^{m} [\hat{y}_j = y_j]$$

where

- $\hat{y}_j = M(\boldsymbol{x}_j)$ is the predicted value from model M of the jth instance $\boldsymbol{x}_j$
- y_j is the ground truth value of the jth instance
- $[P] = \begin{cases} 1 & \text{if } P \text{ is true} \\ 0 & \text{otherwise} \end{cases}$ is the Iverson bracket notation for if/else

Confusion Matrix

Is the student **Alert**?

Student alertness prediction

Inst.	Predicted $\hat{y}$	Actual y	
1	Alert	Alert	TP
2	Alert	Alert	
3	Sleepy	Alert	
4	Sleepy	Alert	FN
5	Sleepy	Alert	
6	Sleepy	Sleepy	
7	Sleepy	Sleepy	TN
8	Sleepy	Sleepy	
9	Sleepy	Sleepy	
10	Alert	Sleepy	FP

Actual Label

Predicted Label	Pos	Neg
Pos	**TP** True Positive	**FP** False Positive
Neg	**FN** False Negative	**TN** True Negative

$$\text{Accuracy} = \frac{\mathrm{TP+TN}}{\mathrm{TP+TN+FP+TN}}$$

Confusion Matrix

Student alertness prediction

Inst.	Predicted $\hat{y}$	Actual y
1	Alert	Alert
2	Alert	Alert
3	Sleepy	Alert
4	Sleepy	Alert
5	Sleepy	Alert
6	Sleepy	Sleepy
7	Sleepy	Sleepy
8	Sleepy	Sleepy
9	Sleepy	Sleepy
10	Alert	Sleepy

Is the student **Alert**?

Predicted Label \ Actual Label	Pos	Neg
Pos	TP True Positive	FP False Positive
Neg	FN False Negative	TN True Negative

Is the student **Sleepy**?

Predicted Label \ Actual Label	Pos	Neg
Pos	TP True Positive	FP False Positive
Neg	FN False Negative	TN True Negative

Which class is Positive? Negative?

You define based on your application

Confusion Matrix

Student alertness prediction

Inst.	Predicted $\hat{y}$	Actual y
1	Alert	Alert
2	Alert	Alert
3	Not	Alert
4	Not	Alert
5	Not	Alert
6	Not	Not
7	Not	Not
8	Not	Not
9	Not	Not
10	Alert	Not

Is the student **Alert**?

		Actual Label	
		Pos	Neg
Predicted Label	Pos	TP True Positive	FP False Positive
	Neg	FN False Negative	TN True Negative

Two types of False mistakes. Which is worse?

- FN: Accuse **alert** students, or
- FP: Neglect **sleepy** students?

Confusion Matrix: which mistake is costlier?

Student alertness prediction

Inst.	Predicted $\hat{y}$	Actual y
1	Alert	Alert
2	Alert	Alert
3	Not	Alert
4	Not	Alert
5	Not	Alert
6	Not	Not
7	Not	Not
8	Not	Not
9	Not	Not
10	Alert	Not

Actual Label

Predicted Label	Pos	Neg
Pos	2 True Positive	1 False Positive
Neg	3 False Negative	4 True Negative

1/10 students predicted
as **Alert** when actually **Sleepy**
=> Few undeserved credits

3/10 students predicted
as **Sleepy** when actually **Alert**
=> Many false accusations

Cost-Sensitive Evaluation Metrics

1. Report **Precision** vs. **Recall**
2. Vary **Prediction Threshold**

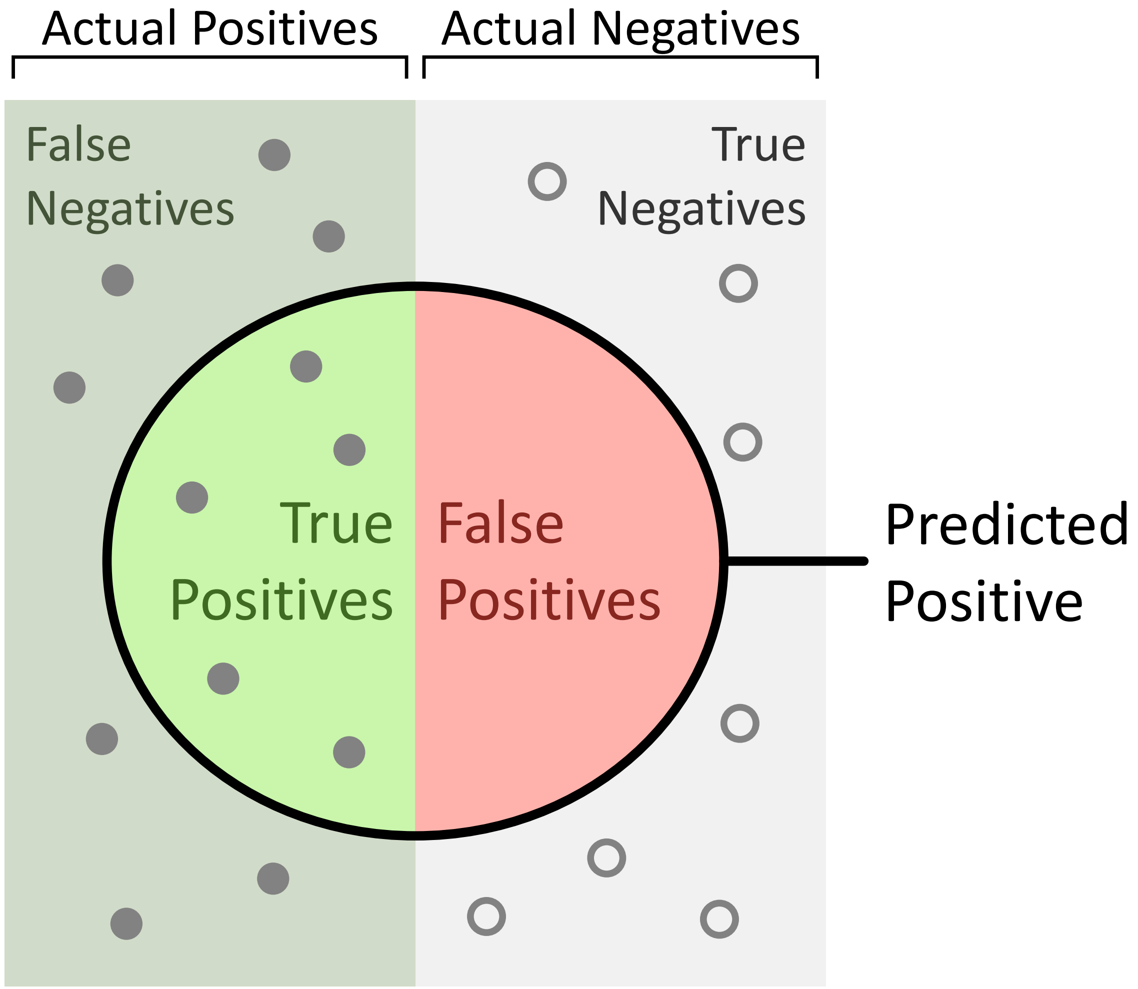

Image credit:
https://en.wikipedia.org/wiki/Precision_and_recall

What fraction of actual positive instances were **recalled**?

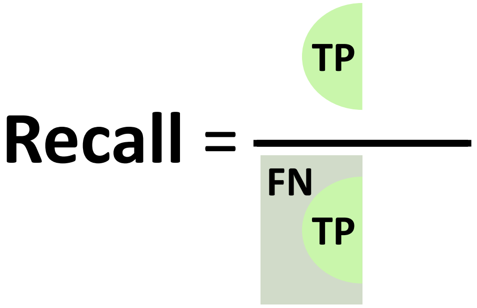

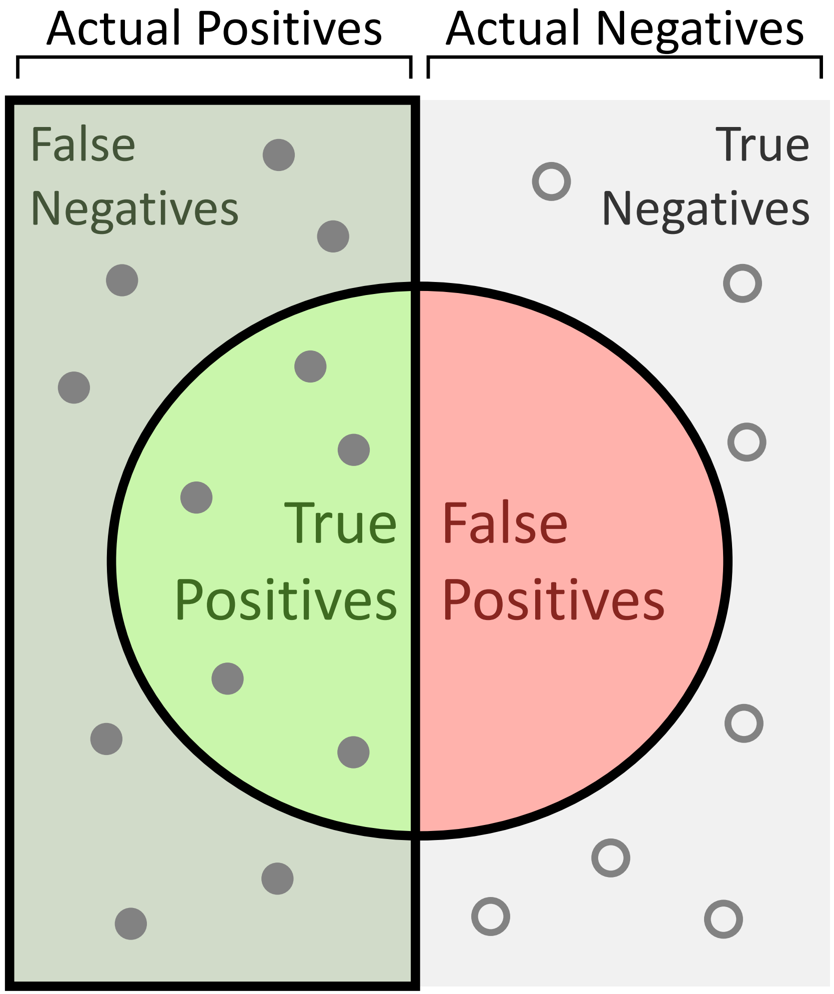

Among positive predictions, how **precisely** were actual positive instances predicted?

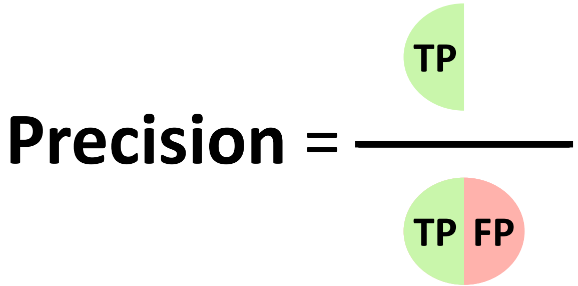

Maximize this
if false negative (FN) is costly.
E.g., cancer prediction,
not music recommendation

Maximize this
if false positive (FP) is costly.
E.g., email spam, satellite launch date prediction.

Image credit:
https://en.wikipedia.org/wiki/Precision_and_recall

Precision and Recall

Inst.	Predicted $\hat{y}$	Actual y
1	Alert	Alert
2	Alert	Alert
3	Not	Alert
4	Not	Alert
5	Not	Alert
6	Not	Not
7	Not	Not
8	Not	Not
9	Not	Not
10	Alert	Not

Actual Label

Predicted Label		Alert	Not	
	Alert	2 True Positive	1 False Positive	3 ∑ Pred. Pos.
	Not	3 False Negative	4 True Negative	7 ∑ Pred. Neg.
		5 ∑ Actual Pos.	5 ∑ Actual Neg.	

Precision
P = TP / (TP+FP)

Recall
R = TP / (TP+FN)

Precision and Recall → F_1 Score

Inst.	Predicted $\hat{y}$	Actual y
1	Alert	Alert
2	Alert	Alert
3	Not	Alert
4	Not	Alert
5	Not	Alert
6	Not	Not
7	Not	Not
8	Not	Not
9	Not	Not
10	Alert	Not

Actual Label

Predicted Label	Alert	Not	
Alert	2 True Positive	1 False Positive	3 ∑ Pred. Pos.
Not	3 False Negative	4 True Negative	7 ∑ Pred. Neg.
	5 ∑ Actual Pos.	5 ∑ Actual Neg.	

Precision
P = TP / (TP+FP)

Recall
R = TP / (TP+FN)

F_1 Score

$$F_1 = \frac{2}{\frac{1}{P} + \frac{1}{R}}$$

F_1 Score: Why not just use simple average?

1. The measure is more **robust** (less sensitive to extreme values)

   Ref: https://stackoverflow.com/a/26360501

2. It considers that the numerators of P and R are the same, so it compares their denominators

$$F_1 = \left(\frac{P^{-1} + R^{-1}}{2}\right)^{-1} = \frac{2}{\frac{1}{P} + \frac{1}{R}} = \frac{2\,TP}{(TP + FP) + (TP + FN)}$$

Other “fairer” metrics that consider true negatives (TN):
Matthews correlation coefficient, Youden’s index, Cohen’s kappa

Imbalanced Data

Balanced

Predicted Label \ Actual Label	Alert	Not	
Alert	2 True Positive	1 False Positive	3 ∑ Pred. Pos.
Not	3 False Negative	4 True Negative	7 ∑ Pred. Neg.
	5 ∑ Actual Pos.	5 ∑ Actual Neg.	

Precision $P = .66$

Accuracy $A = .60$

Recall $R = .40$

F_1 Score $F_1 = .50$

Imbalanced

Predicted Label \ Actual Label	Alert	Not	
Alert	2 True Positive	10 False Positive	12 ∑ Pred. Pos.
Not	3 False Negative	40 True Negative	43 ∑ Pred. Neg.
	5 ∑ Actual Pos.	50 ∑ Actual Neg.	

Precision $P = .17$

Accuracy $A = .76$

Recall $R = .40$

F_1 Score $F_1 = .24$

Imbalanced Data

Predicted Label \ Actual Label	Alert	Not	
Alert	2 True Positive	1 False Positive	3 ∑ Pred. Pos.
Not	3 False Negative	4 True Negative	7 ∑ Pred. Neg.
	5 ∑ Actual Pos.	5 ∑ Actual Neg.	

Precision $P = .66$

Accuracy $A = .60$

Recall $R = .40$

F_1 Score $F_1 = .50$

Predicted Label \ Actual Label	Alert	Not	
Alert	36 True Positive	1 False Positive	37 ∑ Pred. Pos.
Not	14 False Negative	4 True Negative	18 ∑ Pred. Neg.
	50 ∑ Actual Pos.	5 ∑ Actual Neg.	

Precision $P = .97$

Accuracy $A = .73$

Recall $R = .40$

F_1 Score $F_1 = .57$

Cost-Sensitive Evaluation Metrics

1. Report **Precision** vs. **Recall**
2. Vary **Prediction Threshold**

Prediction Confidence

Logistic regression example

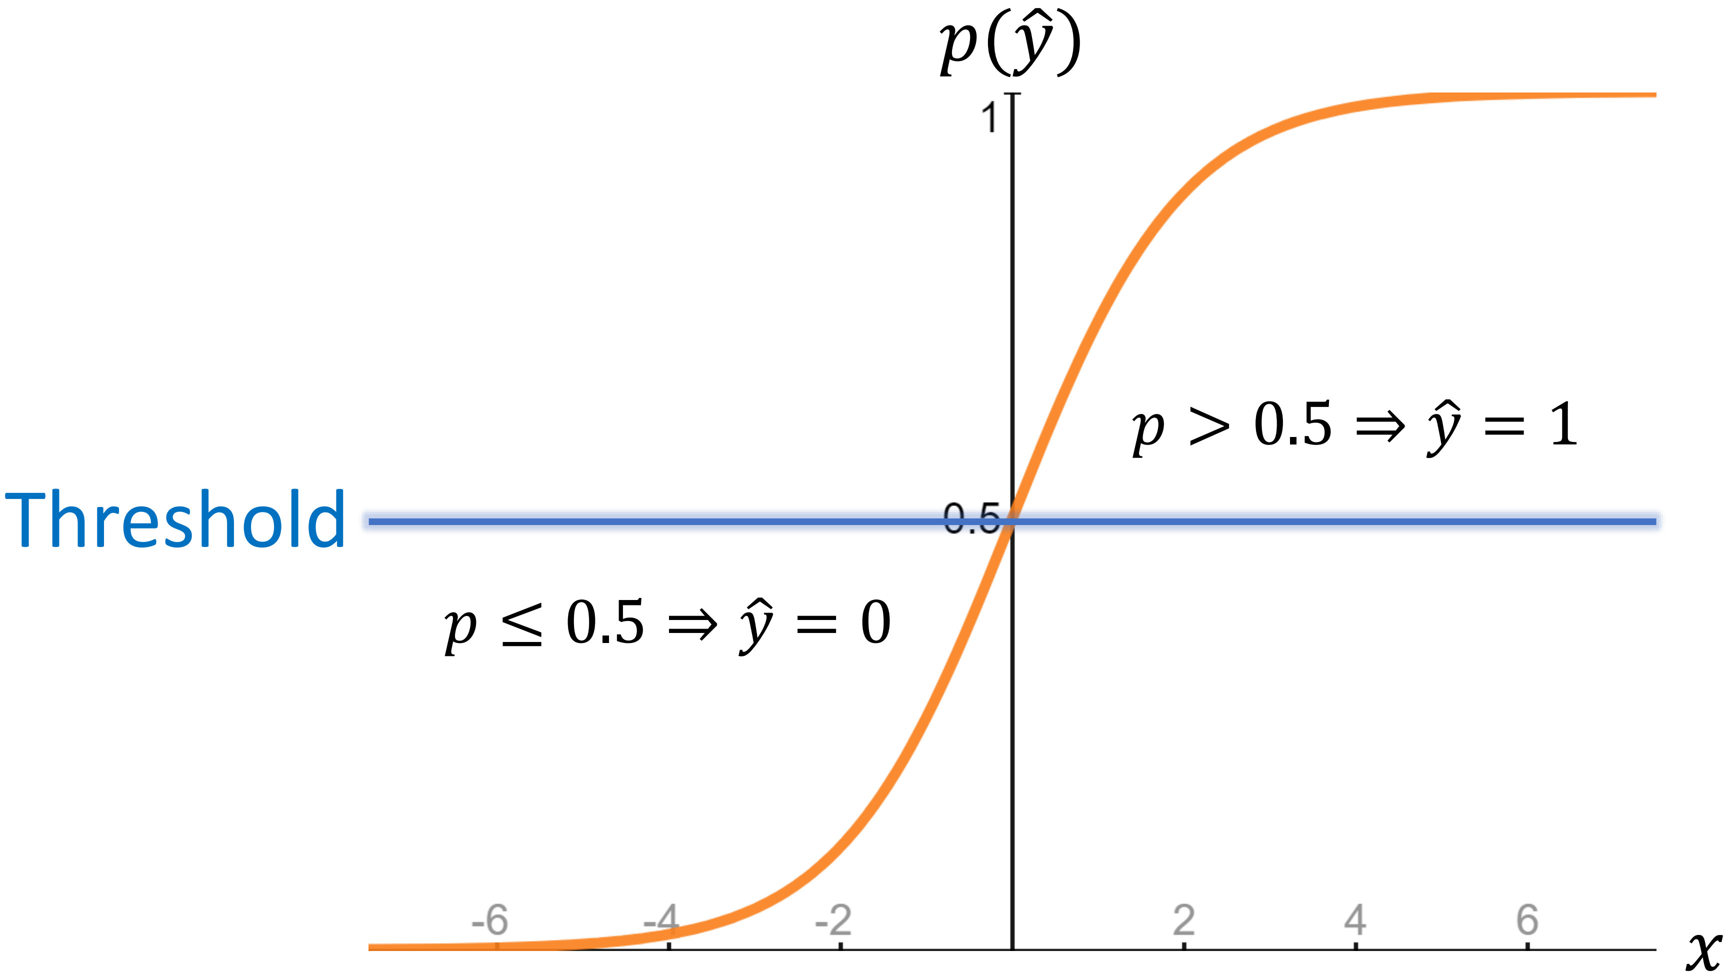

Cost-Sensitive Confusion Matrix (Threshold = **0.5**)

Inst.	**Confidence** $p(\hat{y})$	**Prediction** $\hat{y}$ $p(\hat{y}) > 0.5$	**Actual** y
1	0.9	Alert	Alert
2	0.6	Alert	Alert
3	0.5	Not	Alert
4	0.4	Not	Alert
5	0.3	Not	Alert
6	0.2	Not	Not
7	0.3	Not	Not
8	0.4	Not	Not
9	0.5	Not	Not
10	0.55	Alert	Not

Actual Label

Predicted Label	Alert	Not	
Alert	2 True Positive	1 False Positive	3 ∑ Pred. Pos.
Not	3 False Negative	4 True Negative	7 ∑ Pred. Neg.
	5 ∑ Actual Pos.	5 ∑ Actual Neg.	

True Positive Rate
TPR = TP / (TP+FN)

False Positive Rate
FPR = FP / (FP+TN)

Cost-Sensitive Confusion Matrix (Threshold = **0.5**)

Inst.	Confidence $\mathrm{p}(\hat{y})$	Prediction $\hat{y}$ $\mathrm{p}(\hat{y}) > 0.5$	Actual y
1	0.9	Alert	Alert
2	0.6	Alert	Alert
3	0.5	Not	Alert
4	0.4	Not	Alert
5	0.3	Not	Alert
6	0.2	Not	Not
7	0.3	Not	Not
8	0.4	Not	Not
9	0.5	Not	Not
10	0.55	Alert	Not

Actual Label

Predicted Label	Alert	Not	
Alert	2 True Positive	1 False Positive	3 $\sum$ Pred. Pos.
Not	3 False Negative	4 True Negative	7 $\sum$ Pred. Neg.
	5 $\sum$ Actual Pos.	5 $\sum$ Actual Neg.	

True Positive Rate
TPR = 2/5 = 0.4

False Positive Rate
FPR = 1/5 = 0.2

Cost-Sensitive Confusion Matrix (Threshold = **0.3**)

Inst.	Confidence $p(\hat{y})$	Prediction $\hat{y}$ $p(\hat{y}) > 0.3$	Actual y
1	0.9	Alert	Alert
2	0.6	Alert	Alert
3	0.5	Alert	Alert
4	0.4	Alert	Alert
5	0.3	Not	Alert
6	0.2	Not	Not
7	0.3	Not	Not
8	0.4	Alert	Not
9	0.5	Alert	Not
10	0.55	Alert	Not

		Actual Label		
		Alert	Not	
Predicted Label	Alert	4 True Positive	3 False Positive	7 $\sum$ Pred. Pos.
	Not	1 False Negative	2 True Negative	3 $\sum$ Pred. Neg.
		5 $\sum$ Actual Pos.	5 $\sum$ Actual Neg.	

True Positive Rate
TPR = 4/5 = 0.8

False Positive Rate
FPR = 3/5 = 0.6

Cost-Sensitive Confusion Matrix (Threshold = **0.9**)

Inst.	**Confidence** $\mathrm{p}(\hat{y})$	**Prediction** $\hat{y}$ $\mathrm{p}(\hat{y}) > 0.6$	**Actual** y
1	0.9	Alert	Alert
2	0.6	Not	Alert
3	0.5	Not	Alert
4	0.4	Not	Alert
5	0.3	Not	Alert
6	0.2	Not	Not
7	0.3	Not	Not
8	0.4	Not	Not
9	0.5	Not	Not
10	0.55	Not	Not

Actual Label

Predicted Label	Alert	Not	
Alert	1 True Positive	0 False Positive	1 ∑ Pred. Pos.
Not	4 False Negative	5 True Negative	9 ∑ Pred. Neg.
	5 ∑ Actual Pos.	5 ∑ Actual Neg.	

True Positive Rate
TPR = 1/5 = 0.2

False Positive Rate
FPR = 0/5 = 0.0

Cost-Sensitive Confusion Matrix

…

p > 0.3

Actual Label

		Alert	¬Alert	
Predicted Label	Alert	4 True Positive	3 False Positive	7 ∑ Pred. Pos.
	¬Alert	1 False Negative	2 True Negative	3 ∑ Pred. Neg.
		5 ∑ Actual Pos.	5 ∑ Actual Neg.	

True Positive Rate
TPR = 4/5 = 0.8

False Positive Rate
FPR = 3/5 = 0.6

…

p > 0.5

Actual Label

		Alert	¬Alert	
Predicted Label	Alert	2 True Positive	1 False Positive	3 ∑ Pred. Pos.
	¬Alert	3 False Negative	4 True Negative	7 ∑ Pred. Neg.
		5 ∑ Actual Pos.	5 ∑ Actual Neg.	

True Positive Rate
TPR = 2/5 = 0.4

False Positive Rate
FPR = 1/5 = 0.2

…

p > 0.9

Actual Label

		Alert	¬Alert	
Predicted Label	Alert	1 True Positive	0 False Positive	1 ∑ Pred. Pos.
	¬Alert	4 False Negative	5 True Negative	9 ∑ Pred. Neg.
		5 ∑ Actual Pos.	5 ∑ Actual Neg.	

True Positive Rate
TPR = 1/5 = 0.2

False Positive Rate
FPR = 0/5 = 0.0

…

Confusion matrix depends on **prediction threshold**

Receiver Operator Characteristic (ROC) Curve

Threshold	TPR	FPR
0	1	1
0.1	1.0	1.0
0.2	1.0	0.8
0.3	0.8	0.6
0.4	0.6	0.4
0.5	0.4	0.2
0.6	0.2	0.0
0.7	0.2	0.0
0.8	0.2	0.0
0.9	0.0	0.0
1	0	0

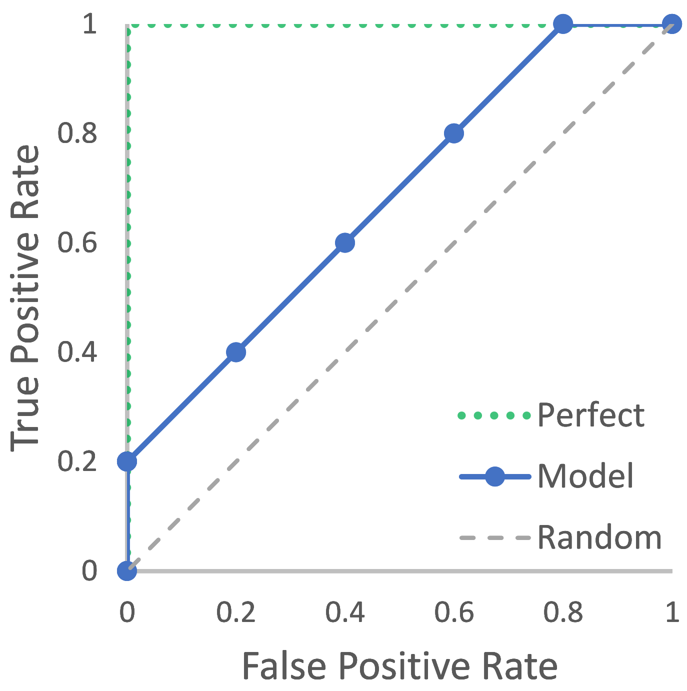

- Diagonal random line indicates 50% chance of correctness.
- If ROC curve is above the random line, model is more accurate than chance.
- Perfect curve has TPR = 1 and FPR = 0 always.

Area Under Curve (AUC) of ROC

Threshold	TPR	FPR
0	1	1
0.1	1.0	1.0
0.2	1.0	0.8
0.3	0.8	0.6
0.4	0.6	0.4
0.5	0.4	0.2
0.6	0.2	0.0
0.7	0.2	0.0
0.8	0.2	0.0
0.9	0.0	0.0
1	0	0

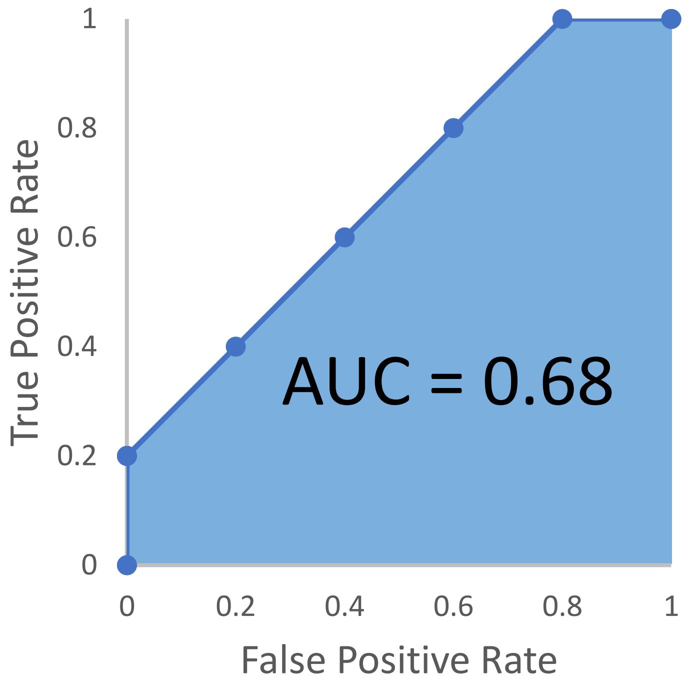

- AUC is a **concise metric** instead of a full figure.
- Concise metrics enable *clearer comparisons*.
- **AUC > 0.5** means the model is better than chance.
- AUC ≈ 1 means model is very accurate.

Area Under Curve (AUC) of ROC (example)

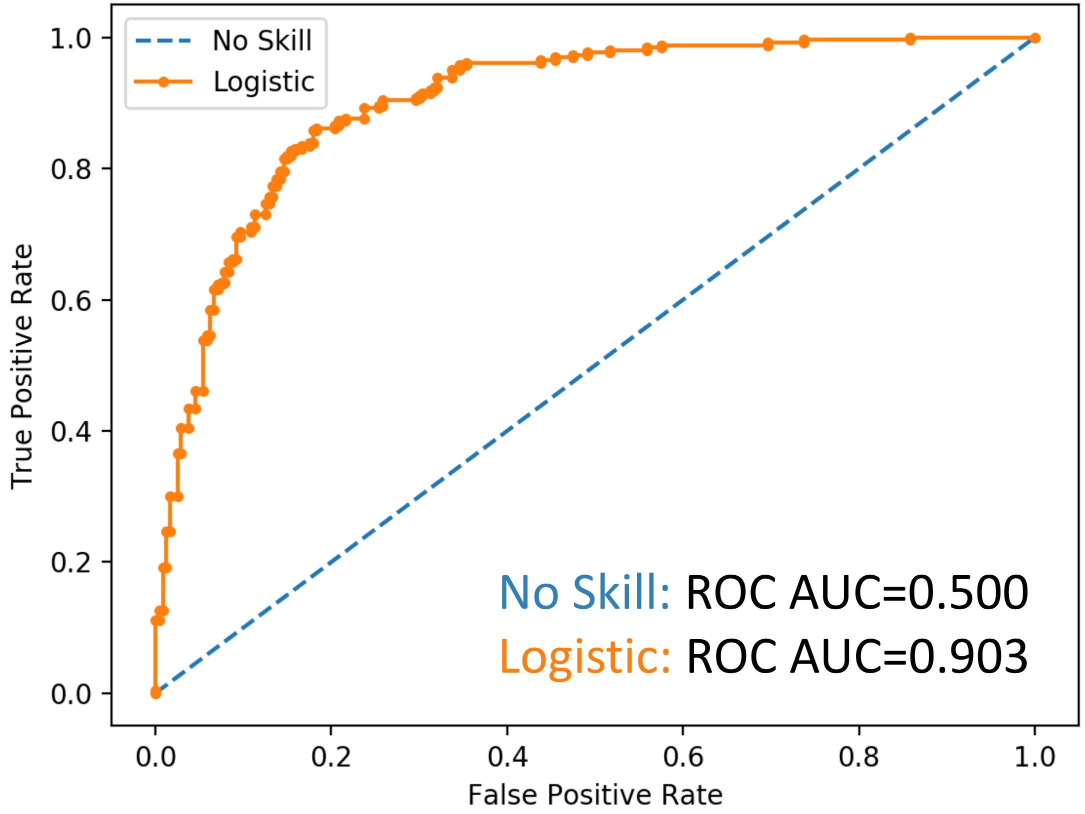

Image credit: https://machinelearningmastery.com/roc-curves-and-precision-recall-curves-for-classification-in-python/


?
Questions!

Classification

$\boxed{\hat{y} \in \{0,1\} \text{ binary}}$

$\boxed{\hat{y} \in \{y_A, y_B, \dots\} \text{ multi-class}}$

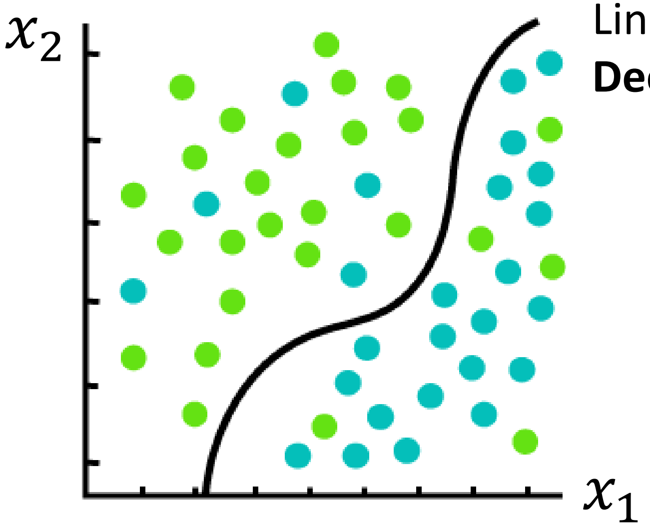

Line is a
Decision Boundary

$$\hat{y} = M(\boldsymbol{x}), \quad \boldsymbol{x} = \vec{x} = (x_1, x_2)^\top$$

Regression

$\hat{y} \in \mathbb{R}$ any real number (scalar)

Line is a
Curve Fit

$$\hat{y} = M(x), \quad x = x_1$$

$$\hat{y} = M(\boldsymbol{x}), \quad \boldsymbol{x} = (x_1, \dots, x_n)^\top$$

Image credit:
https://www.javatpoint.com/regression-vs-classification-in-machine-learning

Confusion Matrix (binary classification)

Inst.	Predicted $\hat{y}$	Actual y
1	Alert	Alert
2	Alert	Alert
3	Sleepy	Alert
4	Sleepy	Alert
5	Sleepy	Alert
6	Sleepy	Sleepy
7	Sleepy	Sleepy
8	Sleepy	Sleepy
9	Sleepy	Sleepy
10	Alert	Sleepy

		Actual Label	
		Alert	Sleepy
Predicted Label	Alert	2	1
	Sleepy	3	4

Confusion Matrix (multiclass classification)

Inst.	Predicted $\hat{y}$	Actual y
1	Alert	Alert
2	Alert	Alert
3	Sleepy	Alert
4	Sleepy	Alert
5	Sleepy	Alert
6	Sleepy	Sleepy
7	Sleepy	Sleepy
8	Sleepy	Sleepy
9	Sleepy	Sleepy
10	Alert	Sleepy
11	Away	Away
12	Away	Alert
…	…	…

Actual Label

Predicted Label	Alert	Sleepy	Away
Alert	3	2	#
Sleepy	2	3	#
Away	#	#	#

Confusion Matrix (multiclass example)

Inst.	Predicted $\hat{y}$	Actual y
1	Alert	Alert
2	Alert	Alert
3	Sleepy	Alert
4	Sleepy	Alert
5	Sleepy	Alert
6	Sleepy	Sleepy
7	Sleepy	Sleepy
8	Sleepy	Sleepy
9	Sleepy	Sleepy
10	Alert	Sleepy
11	Away	Away
12	Away	Alert
…	…	…

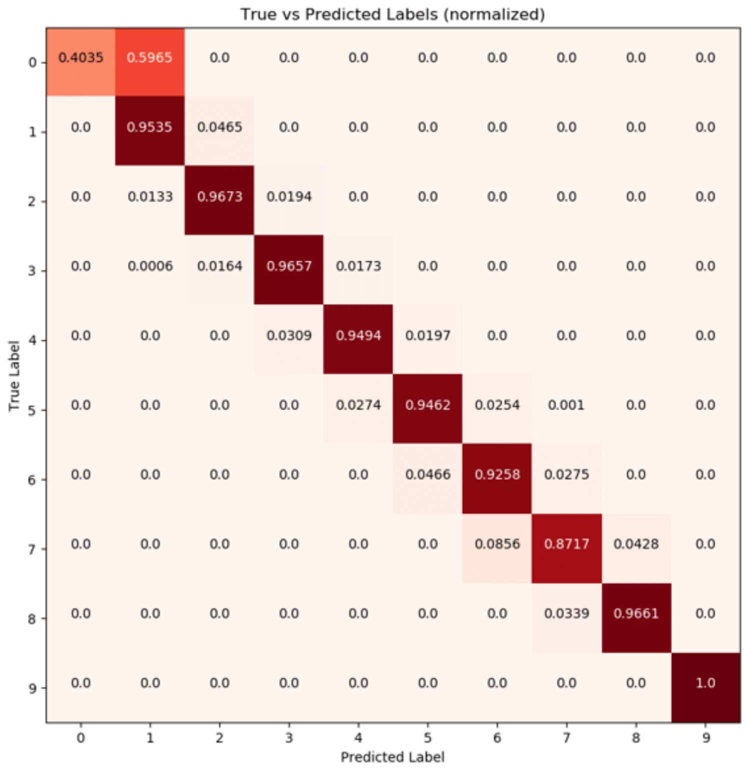

Image Credit: https://www.researchgate.net/figure/Confusion-matrix-for-insample-test-predictions-Notice-the-heavily-diagonally-dominant_fig2_328997503

How to calculate:

- Accuracy
- Precision, Recall, F_1
- AUC?

Confusion Matrix (binary classification)

Inst.	Predicted $\hat{y}$	Actual y
1	Alert	Alert
2	Alert	Alert
3	Sleepy	Alert
4	Sleepy	Alert
5	Sleepy	Alert
6	Sleepy	Sleepy
7	Sleepy	Sleepy
8	Sleepy	Sleepy
9	Sleepy	Sleepy
10	Alert	Sleepy

Actual Label

Predicted Label		Alert	Sleepy
	Alert	2 True Positive	1 False Positive
	Sleepy	3 False Negative	4 True Negative

Confusion Matrix (multiclass classification)

Inst.	Predicted $\hat{y}$	Actual y
1	Alert	Alert
2	Alert	Alert
3	Sleepy	Alert
4	Sleepy	Alert
5	Sleepy	Alert
6	Sleepy	Sleepy
7	Sleepy	Sleepy
8	Sleepy	Sleepy
9	Sleepy	Sleepy
10	Alert	Sleepy
11	Away	Away
12	Away	Alert
…	…	…

Actual Label

Predicted Label	Alert	Sleepy	Away
Alert	2 ~~True Positive~~	1 ~~False Positive~~	#
Sleepy	3 ~~False Negative~~	4 ~~True Negative~~	#
Away	#	#	#

Which class is Positive? Negative?

Confusion Matrix (multiclass classification)

Inst.	Predicted $\hat{y}$	Actual y
1	Alert	Alert
2	Alert	Alert
3	Sleepy	Alert
4	Sleepy	Alert
5	Sleepy	Alert
6	Sleepy	Sleepy
7	Sleepy	Sleepy
8	Sleepy	Sleepy
9	Sleepy	Sleepy
10	Alert	Sleepy
11	Away	Away
12	Away	Alert
…	…	…

Actual Label

Predicted Label	Alert	Not	Not
Alert	TP	FP	FP
Not	FN	TN	TN
Not	FN	TN	TN

Alert class is Positive, others Neg.

Confusion Matrix (multiclass classification)

Inst.	Predicted $\hat{y}$	Actual y
1	Alert	Alert
2	Alert	Alert
3	Sleepy	Alert
4	Sleepy	Alert
5	Sleepy	Alert
6	Sleepy	Sleepy
7	Sleepy	Sleepy
8	Sleepy	Sleepy
9	Sleepy	Sleepy
10	Alert	Sleepy
11	Away	Away
12	Away	Alert
…	…	…

Actual Label

Predicted Label	Not	Sleepy	Not
Not	TN	FN	TN
Sleepy	FP	TP	FP
Not	TN	FN	TN

Sleepy class is Positive, others Neg.

Confusion Matrix (multiclass classification)

Inst.	Predicted $\hat{y}$	Actual y
1	Alert	Alert
2	Alert	Alert
3	Sleepy	Alert
4	Sleepy	Alert
5	Sleepy	Alert
6	Sleepy	Sleepy
7	Sleepy	Sleepy
8	Sleepy	Sleepy
9	Sleepy	Sleepy
10	Alert	Sleepy
11	Away	Away
12	Away	Alert
…	…	…

Actual Label

Predicted Label	Not	Not	Away
Not	TN	TN	FN
Not	TN	TN	FN
Away	FP	FP	TP

Away class is Positive, others Neg.

Multiclass evaluation metrics: **Average?**

Predicted Label	Actual Label →	Alert	Not	Not
	Alert	TP_{c1}	FP_{c1}	FP_{c1}
	Not	FN_{c1}	TN_{c1}	TN_{c1}
	Not	FN_{c1}	TN_{c1}	TN_{c1}

Predicted Label	Actual Label →	Not	Sleepy	Not
	Not	TN_{c2}	FN_{c2}	TN_{c2}
	Sleepy	FP_{c2}	TP_{c2}	FP_{c2}
	Not	TN_{c2}	FN_{c2}	TN_{c2}

Predicted Label	Actual Label →	Not	Not	Away
	Not	TN_{c3}	TN_{c3}	FN_{c3}
	Not	TN_{c3}	TN_{c3}	FN_{c3}
	Away	FP_{c3}	FP_{c3}	TP_{c3}

How to combine?

Multiclass evaluation metrics: **Micro-Average**

Predicted Label \ Actual Label	Alert	Not	Not
Alert	TP_{c1}	FP_{c1}	FP_{c1}
Not	FN_{c1}	TN_{c1}	TN_{c1}
Not	FN_{c1}	TN_{c1}	TN_{c1}

Predicted Label \ Actual Label	Not	Sleepy	Not
Not	TN_{c2}	FN_{c2}	TN_{c2}
Sleepy	FP_{c2}	TP_{c2}	FP_{c2}
Not	TN_{c2}	FN_{c2}	TN_{c2}

Predicted Label \ Actual Label	Not	Not	Away
Not	TN_{c3}	TN_{c3}	FN_{c3}
Not	TN_{c3}	TN_{c3}	FN_{c3}
Away	FP_{c3}	FP_{c3}	TP_{c3}

$$TP_\Sigma = \frac{1}{|C|}\sum_{c\in C} TP_c = (TP_{c1} + TP_{c2} + TP_{c3})/3$$

$$TN_\Sigma = \frac{1}{|C|}\sum_{c\in C} TN_c$$

$$FP_\Sigma = \frac{1}{|C|}\sum_{c\in C} FP_c$$

$$FN_\Sigma = \frac{1}{|C|}\sum_{c\in C} FN_c$$

Predicted Label \ Actual Label	Pos	Neg
Pos	TP_Σ	FP_Σ
Neg	FN_Σ	TN_Σ

$$\boxed{A_\Sigma, P_\Sigma, R_\Sigma, F1_\Sigma, AUC_\Sigma}$$

Multiclass evaluation metrics: **Macro-Average**

Predicted Label \ Actual Label	Alert	Not	Not
Alert	TP_{c1}	FP_{c1}	FP_{c1}
Not	FN_{c1}	TN_{c1}	TN_{c1}
Not	FN_{c1}	TN_{c1}	TN_{c1}

Predicted Label \ Actual Label	Not	Sleepy	Not
Not	TN_{c2}	FN_{c2}	TN_{c2}
Sleepy	FP_{c2}	TP_{c2}	FP_{c2}
Not	TN_{c2}	FN_{c2}	TN_{c2}

Predicted Label \ Actual Label	Not	Not	Away
Not	TN_{c3}	TN_{c3}	FN_{c3}
Not	TN_{c3}	TN_{c3}	FN_{c3}
Away	FP_{c3}	FP_{c3}	TP_{c3}

$$A_{c1}, P_{c1}, R_{c1}, F1_{c1}, AUC_{c1}$$

$$A_{c2}, P_{c2}, R_{c2}, F1_{c2}, AUC_{c2}$$

$$A_{c3}, P_{c3}, R_{c3}, F1_{c3}, AUC_{c3}$$

$$A_{\Sigma}, P_{\Sigma}, R_{\Sigma}, F1_{\Sigma}, AUC_{\Sigma}$$

$$A_{\Sigma} = \frac{1}{|C|}\sum_{c\in C} A_c \qquad P_{\Sigma} = \frac{1}{|C|}\sum_{c\in C} P_c \qquad F1_{\Sigma} = \frac{1}{|C|}\sum_{c\in C} F1_c$$

$$R_{\Sigma} = \frac{1}{|C|}\sum_{c\in C} R_c \qquad AUC_{\Sigma} = \frac{1}{|C|}\sum_{c\in C} AUC_c$$

Micro- vs. Macro Average

Predicted Label \ Actual Label	Alert	Not	Not
Alert	TP_{c1}	FP_{c1}	FP_{c1}
Not	FN_{c1}	TN_{c1}	TN_{c1}
Not	FN_{c1}	TN_{c1}	TN_{c1}

Predicted Label \ Actual Label	Not	Sleepy	Not
Not	TN_{c2}	FN_{c2}	TN_{c2}
Sleepy	FP_{c2}	TP_{c2}	FP_{c2}
Not	TN_{c2}	FN_{c2}	TN_{c2}

Predicted Label \ Actual Label	Not	Not	Away
Not	TN_{c3}	TN_{c3}	FN_{c3}
Not	TN_{c3}	TN_{c3}	FN_{c3}
Away	FP_{c3}	FP_{c3}	TP_{c3}

How to combine?

Micro-Average

Weighs each **instance** equally.
Controls for imbalanced data.
Will have higher precision if data is Imbalanced.

Macro-Average

Weighs each **class** equally.
Use this if you think your real-world test data is balanced.

Imbalanced Classification evaluation with Precision-Recall (PR) Curve AUC

Further reading: https://machinelearningmastery.com/roc-curves-and-precision-recall-curves-for-imbalanced-classification/

Balanced Classes | Imbalanced Classes

ROC Curve

PR Curve

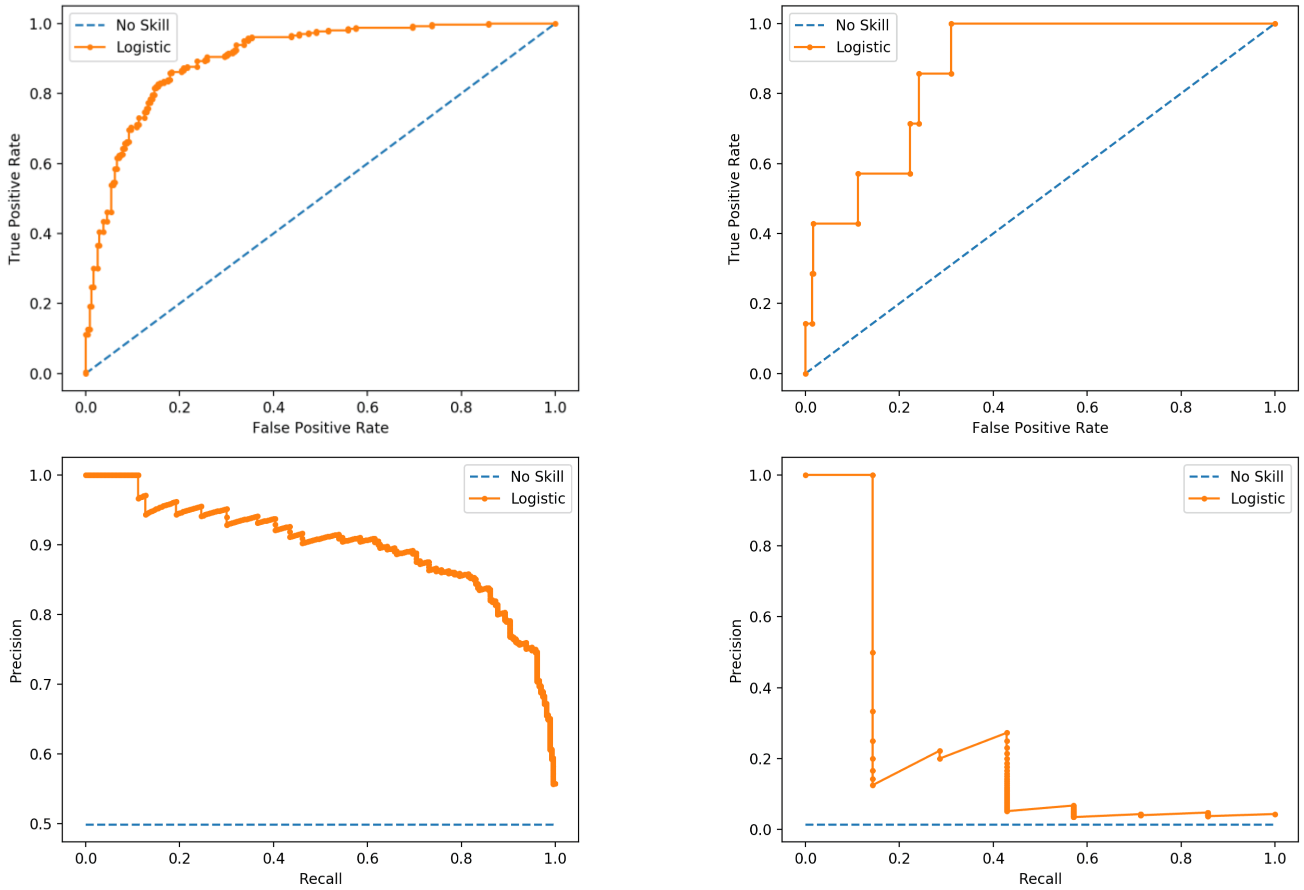

Match appropriate **evaluation metric** to **challenge**

<table>
<tr><th>Challenge</th><th>Evaluation Metric</th></tr>
<tr><td>Imbalanced actual classes</td><td rowspan="3">1️⃣ Accuracy (Emote :one:)
2️⃣ Precision (:two:)
3️⃣ Recall (:three:)
4️⃣ F_1 Score (:four:)
5️⃣ ROC AUC (:five:)
6️⃣ PRC AUC (:six:)
7️⃣ Micro-Average (:seven:)
8️⃣ Macro-Average (:eight:)</td></tr>
<tr><td>Multiclass classification</td></tr>
<tr><td>Cost-dependent classes</td></tr>
</table>

Emote (react) in Slack #general channel <u>one or more options</u> (MRQ) for each challenge

Wrapping Up

What did we learn this week?

1. Recap: Classification vs. Regression
2. Classification Metrics
    1. Accuracy
    2. Confusion Matrix, TP, TN, FP, FN
    3. Precision, Recall, F_1
    4. ROC, AUC
    5. Micro- and Macro-Averaging
    6. PR-AUC (Average Precision)
3. Regression Metrics [W08a]

Performance Metrics

Regression Evaluation Metrics

NUS National University of Singapore | Department of Computer Science School of Computing

Next week: Data Preparation for ML

Image credit:
https://img2.thejournal.ie/article/5047666/river?version=5047733&width=1340

W08 Pre-Lecture Task (due before next Mon)

Read

1. Discover Feature Engineering, How to Engineer Features and How to Get Good at It by Jason Brownlee
2. 8 Tactics to Combat Imbalanced Classes in Your Machine Learning Dataset by Jason Brownlee

Task

1. Identify cases of **bad data** in machine learning
2. Propose **mitigation strategies**

   Tip: you can your own projects too; you don't have to be correct
3. Post a 1–2 sentence answer to the topic in your tutorial group: #tg-xx