

# Data Processing

# 8


# A

**CS 3244**  
**Machine Learning**

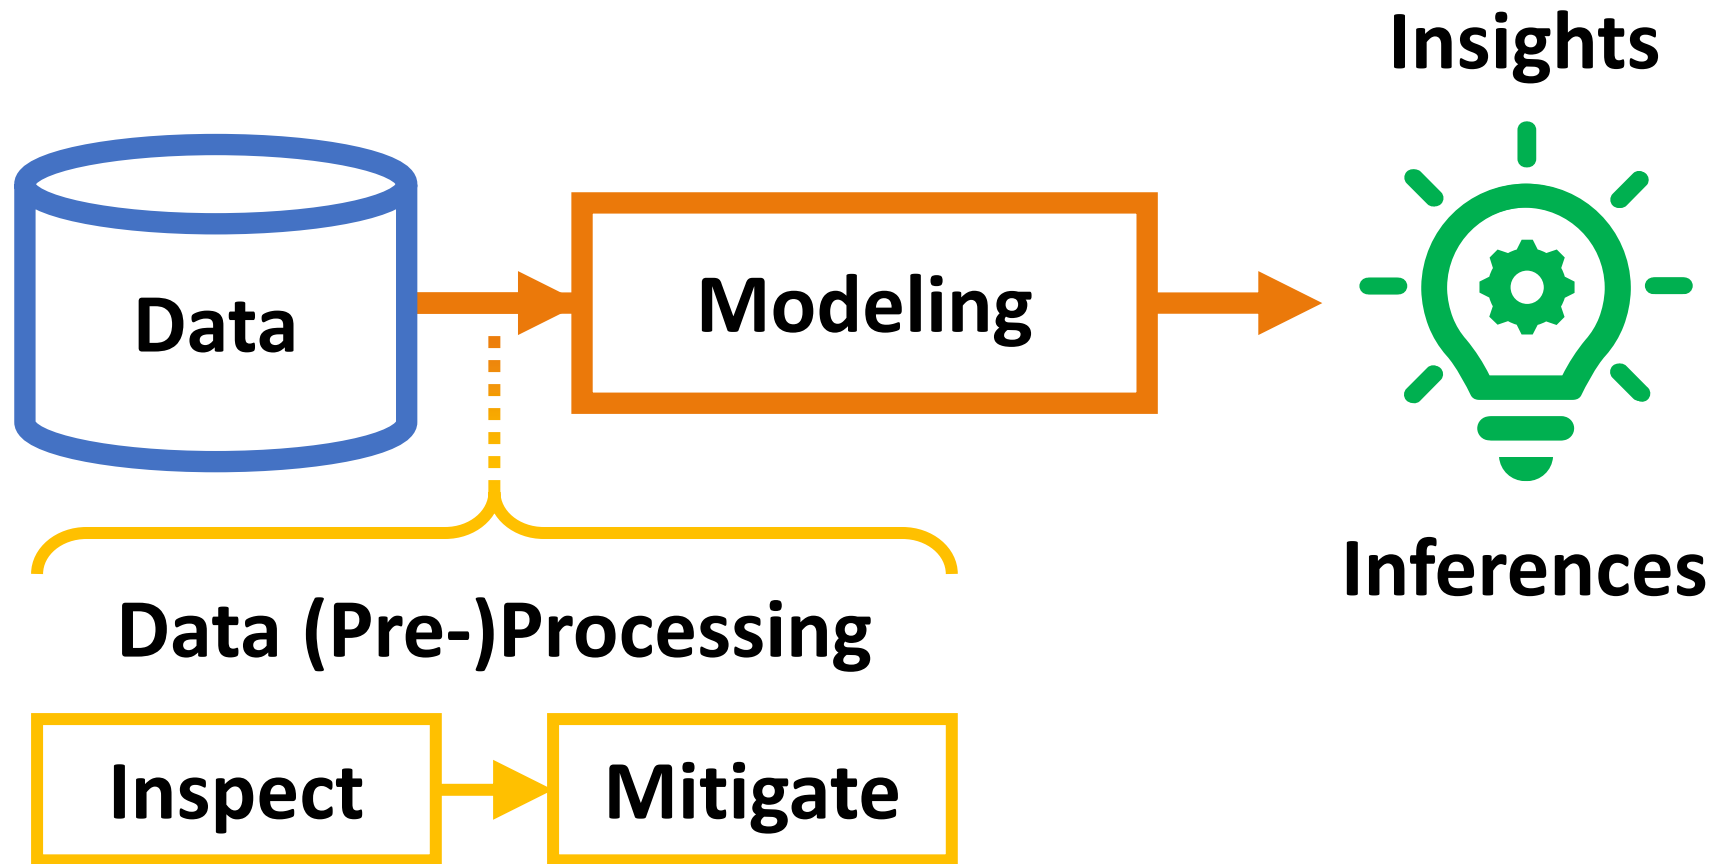


**Computing**

# Machine Learning Pipeline



# Machine Learning Pipeline



# W08 Pre-Lecture Task

## Read

1. [Discover Feature Engineering, How to Engineer Features and How to Get Good at It](#) by Jason Brownlee
2. [8 Tactics to Combat Imbalanced Classes in Your Machine Learning Dataset](#) by Jason Brownlee

## Task

1. Identify cases of **bad data** in machine learning
2. Propose **mitigation strategies**

**Tip:** you can your own projects too; you don't have to be correct

3. Post a 1–2 sentence answer to the topic in your tutorial group: **#tg-xx**

1. Identify cases of **bad data** in machine learning
2. Propose **mitigation strategies**

1 Bad data could include statistical noise and errors. To mitigate this, do data cleaning.

2 Bad data could be when there are too many irrelevant features, which can be mitigated by selecting only the most important features, or by feature extraction where several features are combined into a single feature that is more relevant.

3 High dimensional features

- a. Since the dataset is made up of raw images, each image is represented by a  $256 \times 256 \times 3$  array of numbers which, if simply flattened without much feature engineering, will result in 196608 dimensions. This will be an issue as the model will not only suffer from the curse of high dimensionality, it will also take extremely long to train.
- b. **Mitigation:** I plan to find different ways to reduce the dimensionality of the dataset during the pre-processing phase and choose the one that gives the best results.

4 To mitigate the problem of imbalanced data, we could try to generate synthetic samples of the minority class. We could also try to change our performance metric to take into account the imbalance of our dataset. For example, the Cohen's Kappa can provide a classification accuracy normalised by the imbalance of the classes in the data. To reduce the noise in the dataset, we could resample or collect more data.

There are many different cases of bad data such as imbalanced data, missing data or error/mistakes in the collected data etc.

Mitigation strategies for such cases of bad data could include resampling the data, collecting more data or in some cases even just dropping the erroneous/missing data

<https://www.reuters.com/article/us-amazon-com-jobs-automation-insight-idUSKCN1MK08G>

Here is an example of how imbalanced data can cause real-world harms. Amazon fed in the resumes of people who were hired at Amazon to create a model to screen applicants. However, since the men outnumbered the women in the training data, the model learned a hypothesis that was skewed on favour of male applicants. This shows how if one is not careful with the composition of the training data in a deployed model, it can lead to real-world harms. Possible methods to combat this could be to feed the model fictional female candidates to try to make the classes more even, and thus make the predictions less skewed.

U.S.

Amazon scraps secret AI recruiting tool that showed bias against women (43 kB)



Issues identified:

1. Erroneous data
2. Irrelevant data
3. High dimensionality
4. Imbalanced data
5. Missing data

5

One common case of bad data is data with a lot of missing values. Some ways to mitigate this are to delete the data (either likewise or pairwise deletion) or to impute the missing values (e.g. replace the data with mean, median, or mode, or do regression to get the "predicted" values)

# Week 08: Learning Outcomes

## Data Issues

1. Linear Separability
2. Curse of Dimensionality
3. Imbalanced Data

## Issue Template

1. **What** is the issue?
2. **Why** is it a problem?
3. **When** would it happen?
4. **How** to **check** for it?
5. **How** to **mitigate** it?

For each issue, which of the following techniques can:

- 1) **Check** for the issue?
- 2) **Mitigate** the issue?

Issue	Check	Mitigate
Linear Separability		1 Feature Engineering
Curse of Dimensionality		2 Feature Extraction (extract new features)
		3 Information Gain
		4 Linear Discriminant Analysis (LDA)
Imbalanced Data		5 Principle Components Analysis (PCA)
		6 SMOTE
		7 Support Vector Machine
		8 Visualize Histogram
		9 Visualize Scatterplot

Emote (react) in Slack [#general](#) channel one or more options (MRQ) for each issue

For each issue, which of the following techniques can:  
1) **Check** for the issue?  
2) **Mitigate** the issue?

Issue	Check	Mitigate
Linear Separability	<div><div>9</div> Visualize Scatterplot</div> <div><div>7</div> Support Vector Machine</div> <div><div>4</div> Check Basis Vectors (with</div> <div><div>5</div> LDA, PCA)</div>	<div><div>1</div> Feature Engineering</div> <div><div>2</div> Feature Extraction</div> <div><div>4</div> Matrix Factorization (with</div> <div><div>5</div> LDA, PCA)</div>
Curse of Dimensionality	<div><div>8</div> Visualize Histogram (of distances)</div>	<div><div>3</div> Feature Selection (using Information Gain)</div> <div><div>4</div> Dimensionality Reduction</div> <div><div>5</div> (with LDA, PCA)</div>
Imbalanced Data	<div><div>8</div> Visualize Histogram</div>	<div><div>6</div> SMOTE</div>






# Linear Separability


# Linearly Separable?

Yes



$$\mathbf{x} = \begin{pmatrix} x_1 \\ x_2 \end{pmatrix}$$

**Not** without  
data processing



**How** to make linearly separable?


$$\mathbf{x}' = \begin{pmatrix} (x_1 - \bar{x}_1)^2 \\ (x_2 - \bar{x}_2)^2 \end{pmatrix} = (\mathbf{x} - \bar{\mathbf{x}})^\top (\mathbf{x} - \bar{\mathbf{x}})$$

**Feature Engineering!**


Which of the following is:

- 1. Is Linearly separable (📏:straight\_ruler:)?
- 2. Can it be made Linearly Separable (🌀:curly\_loop:)? How? (Write in **thread**)


a)




b)



c)




d)




Which of the following is:  
1. Is Linearly separable (📏:straight\_ruler:)?  
2. Can it be made Linearly Separable (🌀:curly\_loop:)? How? (Write in thread)


a) 🌀📏




b) 🌀



c) 🌀



d) 🌀



# Issue: Linear Separability

## 1. **What** is the issue?

1. Many models assume that data features are **linearly** separable
2. Does your data satisfy this **assumption**?

## 2. **Why** is it a problem?

1. Irrelevant features will be **uninformative** to train the model to discriminate between prediction labels
2. If features are not linearly separable, you **cannot** learn a good **linear model**
3. Need to use more complex models

## 3. **When** would it happen?

1. Most of the time, for “fresh” unprocessed data.
2. Especially for unstructured (non-tabular) data, e.g., images, time, text

# Issue: Linear Separability

## 4. How to check for it?

### 1. Visualize

- 2D: **Scatterplot** of  $x_1$  by  $x_2$  graph
- >2D: **Scatterplot Matrix**
- 500 dimensions?

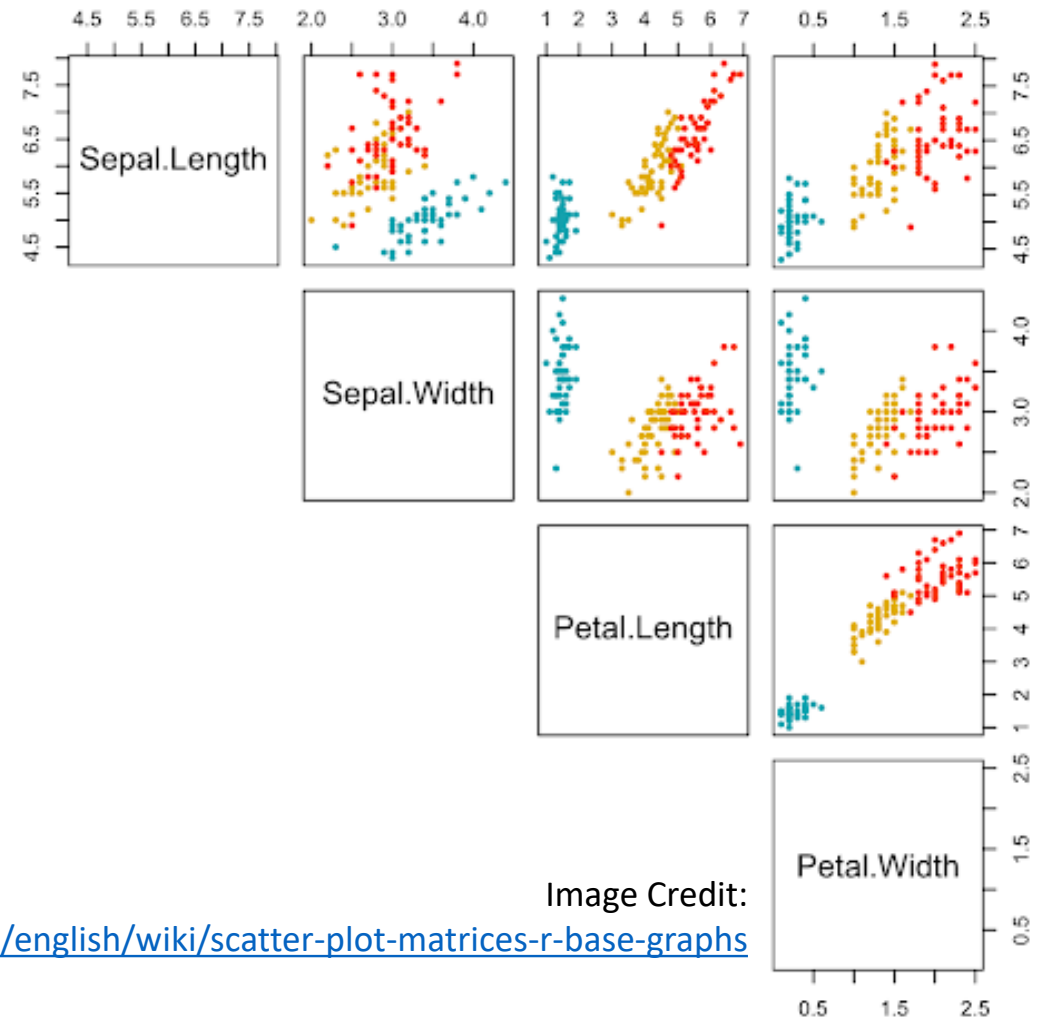


Image Credit:

<http://www.sthda.com/english/wiki/scatter-plot-matrices-r-base-graphs>



# Issue: Linear Separability

## 4. How to check for it?

1. Visualize
2. Computational metrics
  1. Linear SVM [W04b]




Image Credit:

<http://www.sthda.com/english/wiki/scatter-plot-matrices-r-base-graphs>

# Cost Function w Slack Variables

Margin violation:  $y^{(*)}(\boldsymbol{\theta}^\top \mathbf{x}^{(*)} + b) \geq 1$  fails

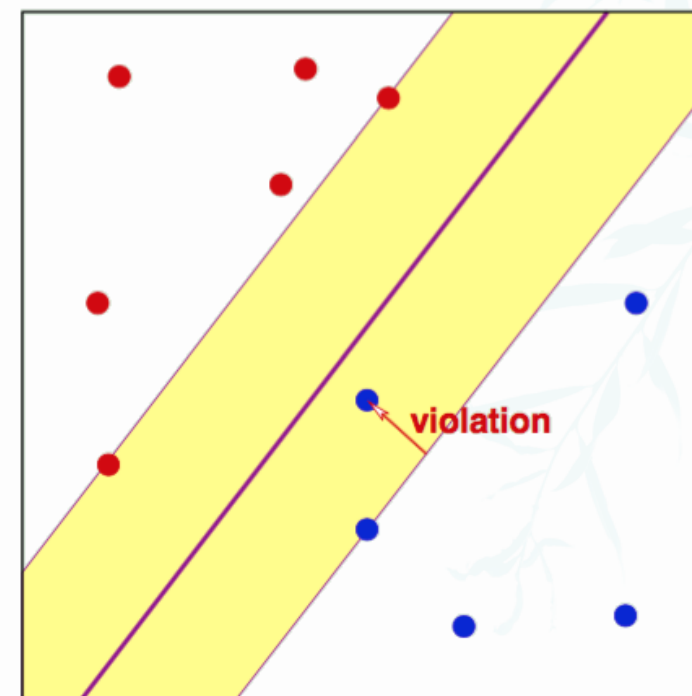
Quantify this:

$$y^{(*)}(\boldsymbol{\theta}^\top \mathbf{x}^{(*)} + b) \geq 1 - \xi^{(*)}$$

where  $\xi^{(*)} \geq 0$

Slack variable: Soft error  
on  $(x^{(*)}, y^{(*)})$


Total violation:  $\sum_{j=1}^m \xi^{(j)}$





# Testing Linear Separability with Linear Soft-Margin SVM

- Each  $\xi^{(j)}$  is the **distance** that the misclassified point  $j$  is from its correct margin
- Total violation:  $\sum_{j=1}^m \xi^{(j)}$
- Calculating the **total violation** indicates how **linearly separable** the data is in terms of its features
- Higher **violation** => **Less** linearly separable



# Issue: Linear Separability

## 4. How to check for it?

1. Visualize
2. Computational metrics
  1. Linear SVM [W04b]
  2. Reduce dimensions (LDA, PCA), then check separability (for separation by “diagonal planes”)
  3. Others: Linear programming, Convex Hulls

Only these are **examinable**



# Issue: Linear Separability

## 5. How to mitigate it?

- Find useful features
  - Feature extraction (collect new features of your data)
- Transformation of features
  - Feature Engineering (e.g.,  $x \rightarrow x^2$ )
  - Change Basis Vectors (e.g., PCA, LDA)
  - Kernel trick (e.g., for kernel SVM [W04b])
  - Feature Learning (e.g., Neural Networks [W09/10])

# Issue: Linear Separability

## 5. How to mitigate it?

- Find useful features
  - Feature extraction (collect new features of your data)
- Transformation of features
  - Feature Engineering (e.g.,  $x \rightarrow x^2$ )
  - **Change Basis Vectors (e.g., PCA, LDA)**
  - Kernel trick (e.g., for kernel SVM [W04b])
  - Feature Learning (e.g., Neural Networks [W09/10])

# Vector Spaces and Basis Vectors




Image Credit: <https://nirpyresearch.com/classification-nir-spectra-linear-discriminant-analysis-python/>

# Vector Spaces and Basis Vectors





Image Credit: <https://nirpyresearch.com/classification-nir-spectra-linear-discriminant-analysis-python/>

# Principal Component Analysis (PCA)

What axis best  
**describes the  
variation** in  
the data?



$$\begin{pmatrix} x_1 \\ x_2 \end{pmatrix} \xrightarrow[\text{PCA Projection}]{s_1 > s_2} \begin{pmatrix} PC_1 \\ PC_2 \end{pmatrix}$$

$$\begin{pmatrix} PC_1 \\ PC_2 \end{pmatrix} \xrightarrow[\text{Reduce Dimensions}]{\text{}} (PC_1)$$


Further reading:

[PCA 1: the basics - simply explained](#) by [TileStats](#),

[StatQuest: Principal Component Analysis \(PCA\), Step-by-Step](#) by [StatQuest with Josh Starmer](#)

Image Credit: <https://nirpyresearch.com/classification-nir-spectra-linear-discriminant-analysis-python/>

# Principal Component Analysis (PCA)



$$\begin{pmatrix} x_1 \\ x_2 \end{pmatrix} \xrightarrow[\text{PCA Projection}]{s_1 > s_2} \begin{pmatrix} PC_1 \\ PC_2 \end{pmatrix}$$

$$\begin{pmatrix} PC_1 \\ PC_2 \end{pmatrix} \xrightarrow[\text{Reduce Dimensions}]{\text{}} (PC_1)$$

Image Credit: <https://ekamperi.github.io/mathematics/2021/02/23/pca-limitations.html>




# Linear Discriminant Analysis (LDA)

What axis best  
**distinguishes**  
**classes** in the  
data?


Maximize

$$F = \frac{|\mu_r - \mu_b|^2}{s_r^2 + s_b^2}$$



Further reading: [Linear discriminant analysis \(LDA\) - simply explained](#) by [TileStats](#),  
[StatQuest: Linear Discriminant Analysis \(LDA\) clearly explained](#) by [StatQuest with Josh Starmer](#)  
Image Credit: <https://nirpyresearch.com/classification-nir-spectra-linear-discriminant-analysis-python/>

# Linear Discriminant Analysis (LDA)




Further reading: [Linear discriminant analysis \(LDA\) - simply explained](#) by [TileStats](#),  
[StatQuest: Linear Discriminant Analysis \(LDA\) clearly explained](#) by [StatQuest with Josh Starmer](#)  
Image Credit: <https://nirpyresearch.com/classification-nir-spectra-linear-discriminant-analysis-python/>

# PCA and LDA

## PCA


Maximize Data Variance



Good for  
**dimensionality reduction**  
for supervised regression

## LDA

Maximize Class Separation



Good for  
**dimensionality reduction**  
for supervised classification  
and unsupervised learning

## Steps

- All axes are orthogonal (independent)
- 1. **Identify** basis vectors
- 2. **Rank** basis vectors by importance
- 3. **Truncate** selection of basis vectors
  - Keeps more important features
  - Performs dimensionality reduction



*Questions!*





# Curse of Dimensionality



# Sparsity with high dimensions




$$m = 5$$

$$n = 1$$




$$m = 25$$

$$n = 2$$



$$m = 125$$


$$n = 3$$



Sparsity problem: maintaining density of samples depends on exponential growth of the data

# Curse of Dimensionality

In high dimensional space, most points are nearly the same distance away.  
The result: learners that depend on distance break down in high dimensions.



<https://stats.stackexchange.com/questions/451027/mathematical-demonstration-of-the-distance-concentration-in-high-dimensions>

# Issue: Curse of Dimensionality

## 1. What is the issue?

1. Too **many features**; many more features than instances

## 2. Why is it a problem?

1. Data **too sparse** to inform about true decision boundary (for classification)  
=> Too easy to fit a model on sparse training data => **Overfitting**
2. Distances are **too similar** (bad for kNN [W02], clustering [W11])

## 3. When would it happen?

1. Extracted more features than data instances (i.e.,  $n \gtrsim m$ )
2. Unstructured data (e.g., features as image pixels, sensor readings)



# Issue: Curse of Dimensionality

## 4. How to check for it?

- Visualize histogram of **distances** (check for **variance**  $\sigma^2$ )

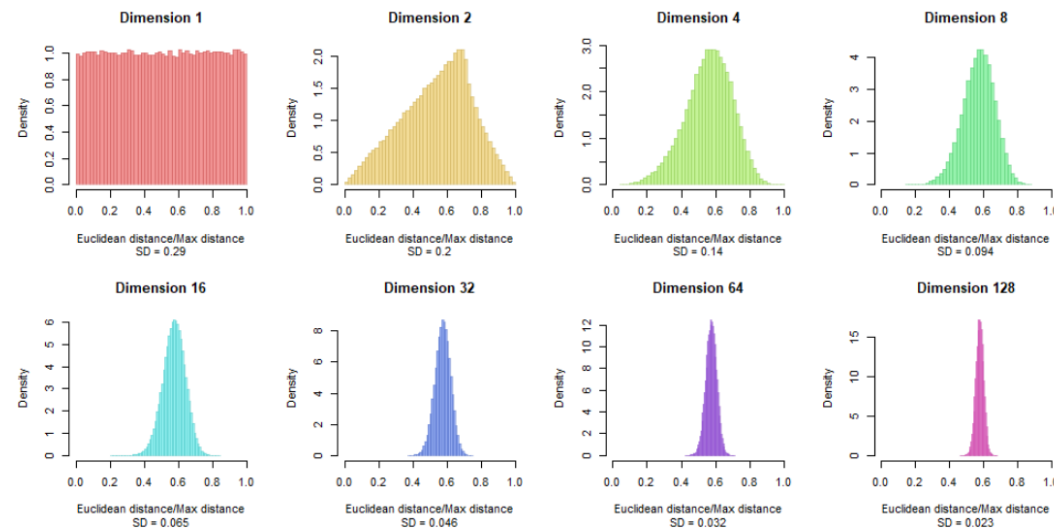


Image Credit: <https://www.mygreatlearning.com/blog/understanding-curse-of-dimensionality/>

Generally **tedious** to analyze this; just aim for:  $n < m/5$

# Issue: Curse of Dimensionality

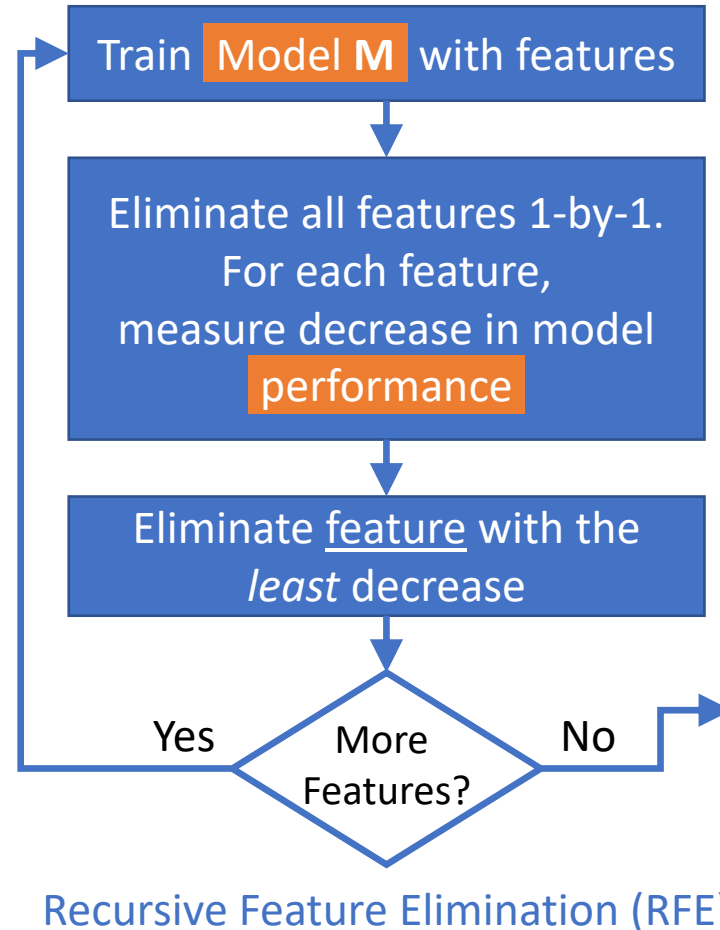
## 5. How to **mitigate** it?

- Feature Selection
  - Wrapper methods
  - Filter methods

# Issue: Curse of Dimensionality

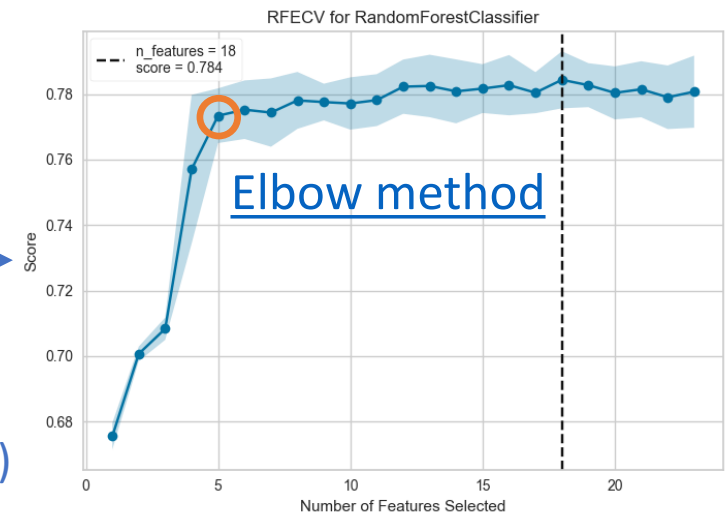
## 5. How to mitigate it?

- Feature Selection
  - **Wrapper methods** (e.g., [RFE](#))
  - Filter methods



Recursive Feature Elimination (RFE)

Image Credit: [https://www.scikit-yb.org/en/latest/api/model\\_selection/rfecv.html](https://www.scikit-yb.org/en/latest/api/model_selection/rfecv.html)



# Issue: Curse of Dimensionality

## 5. How to mitigate it?

- Feature Selection
  - Wrapper methods
  - Filter methods
    - Mutual Information = Information Gain [W03b]
    - Correlation

Recap W03a (slides 22-28)

### Information gain



A chosen feature  $x_i$  divides the example set  $S$  into subsets  $S_1, S_2, \dots, S_c$  according to the  $C_i$  distinct values for  $x_i$ .

The entropy then reduces to the entropy of the subsets  $S_1, S_2, \dots, S_c$ :

$$\text{remainder}(S, x_i) = \sum_{j=1}^{C_i} \frac{|S_j|}{|S|} H(S_j)$$

**Information Gain** (IG; "reduction in entropy") from knowing the value of  $x_i$ .  
Choose the attribute with the largest IG:

$$\text{IG}(S, x_i) = H(S) - \text{remainder}(S, x_i)$$

For the training set at the root,  
 $p = n = 6, H\left(\frac{6}{12}, \frac{6}{12}\right) = 1$  bit.



Consider the attributes **Patrons** and **Type**:

$$\text{IG}(\text{Patrons}) = 1 - \left[ \frac{2}{12} H(0,1) + \frac{4}{12} H(1,0) + \frac{6}{12} H\left(\frac{2}{6}, \frac{4}{6}\right) \right] = 0.541 \text{ bits}$$

$$\text{IG}(\text{Type}) = 1 - \left[ \frac{2}{12} H\left(\frac{1}{2}, \frac{1}{2}\right) + \frac{2}{12} H\left(\frac{1}{2}, \frac{1}{2}\right) + \frac{4}{12} H\left(\frac{2}{4}, \frac{2}{4}\right) + \frac{4}{12} H\left(\frac{2}{4}, \frac{2}{4}\right) \right] = 0 \text{ bits}$$

**Patrons** has the highest IG, and so is chosen by DTL as the root.

Further Reading: <https://machinelearningmastery.com/feature-selection-with-real-and-categorical-data/>,  
<https://towardsdatascience.com/feature-selection-for-machine-learning-3-categories-and-12-methods-6a4403f86543>


# Issue: Curse of Dimensionality

## 5. How to mitigate it?

- Feature Selection
  - Wrapper methods
  - **Filter methods**
    - Mutual Information
    - Correlation

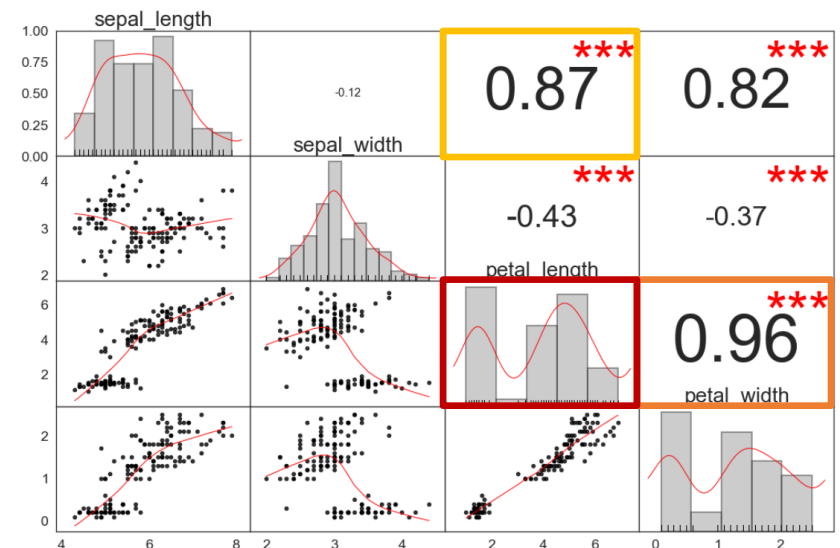
`petal_length` is highly correlated to `petal_width` and `sepal_length`  
=> Should **remove**, due to redundancy

Pearson Correlation Coefficient



Higher **magnitude** => more **correlated**.  $r > 0.7$  is very **high**.  
Further reading: [https://en.wikipedia.org/wiki/Pearson\\_correlation\\_coefficient](https://en.wikipedia.org/wiki/Pearson_correlation_coefficient)

Pearson Correlation Coefficients  
for Iris (flower) dataset



# Issue: Curse of Dimensionality

## 5. How to mitigate it?

- Feature Selection
- Dimensionality Reduction
  - Linear Matrix Factorization (e.g., PCA, LDA)
  - Non-linear Manifold Learning (e.g., SOM, MDS, t-SNE, UMAP)
  - Deep Auto-Encoders

Only these are **examinable**

Further reading:

<https://machinelearningmastery.com/dimensionality-reduction-for-machine-learning/>

# Benefits of Feature Selection


- Avoid **Curse of Dimensionality**
- **Faster** model training (optimizing fewer parameters on fewer features)
- Fewer features to read => **easier to interpret**



# Imbalanced Data

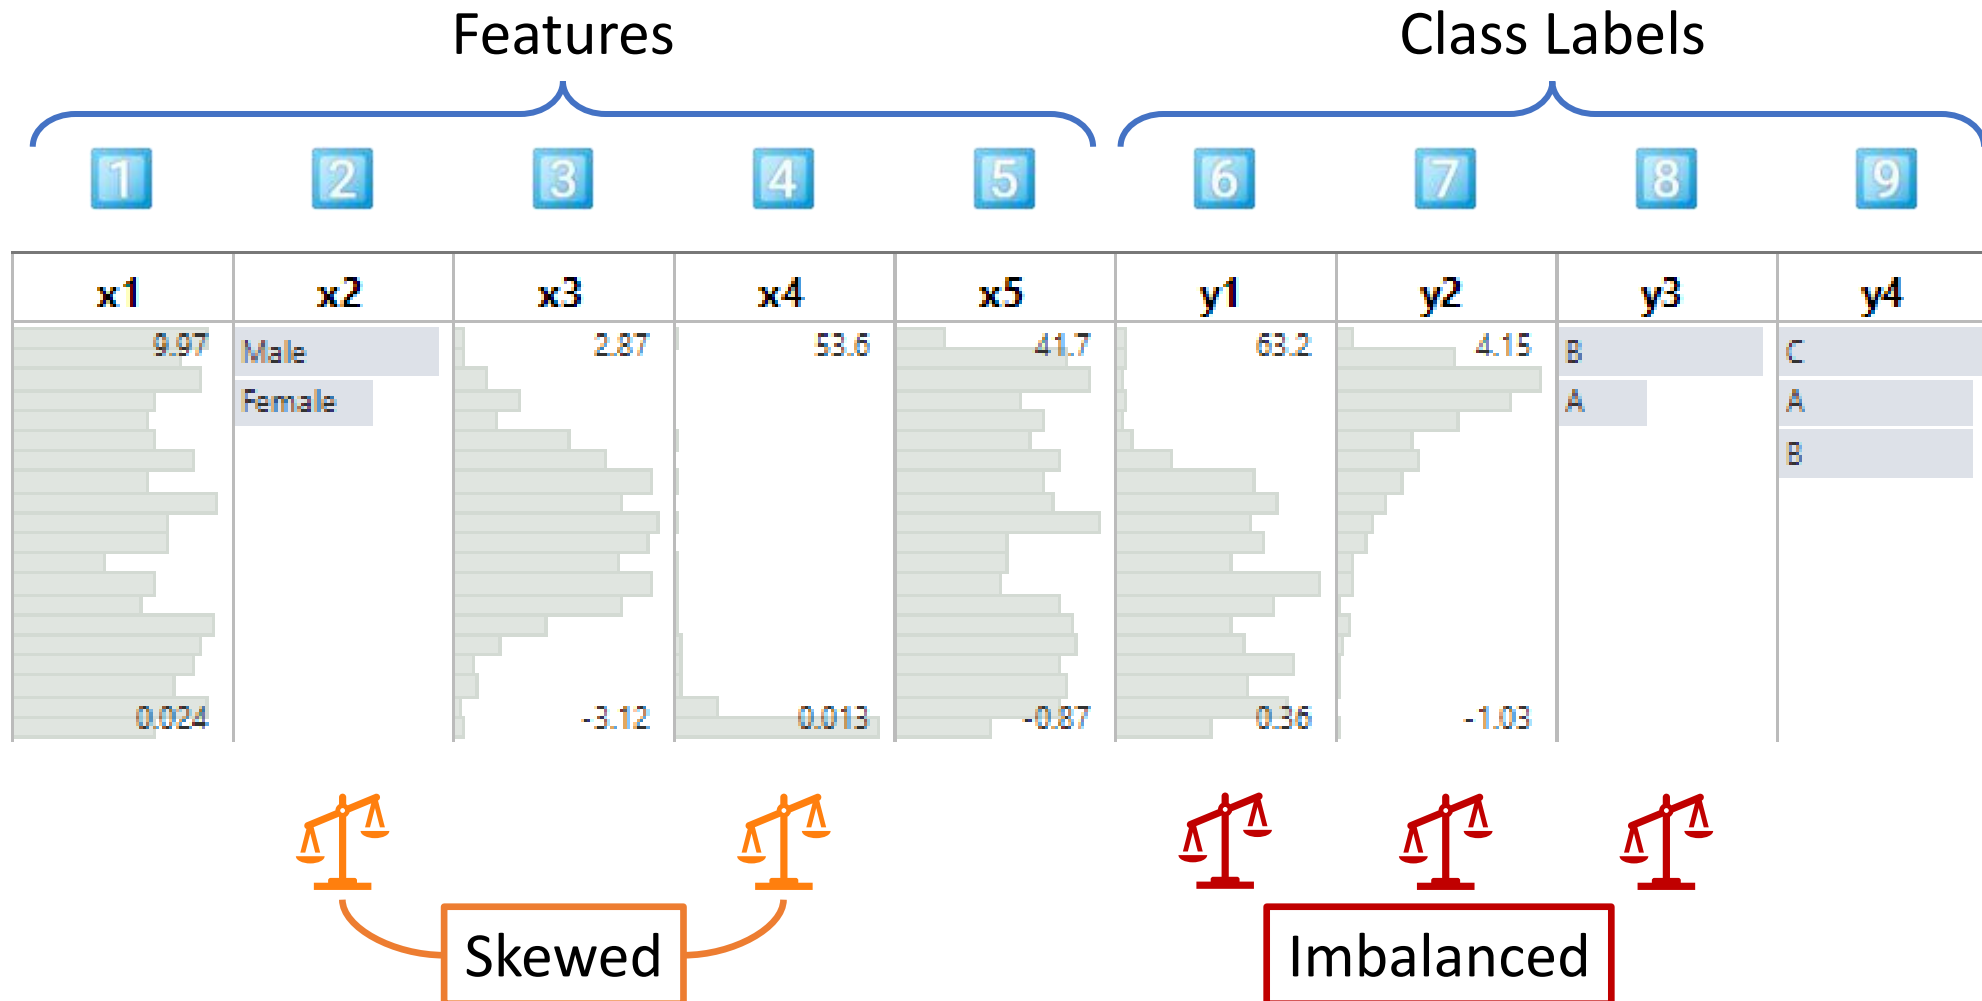


Which of the following variables (columns) are **imbalanced**?



Emote (react) in Slack [#general](#) channel one or more options (MRQ)

Which of the following variables (columns) are imbalanced?



# Issue: Imbalanced Data

## 1. **What** is the issue?

1. Values **not evenly distributed** in feature
2. Data may be skewed

## 2. **Why** is it a problem?

1. Evaluation metrics become **misleading** to interpret [W07b]
2. Models **overfit** to majority class


## 3. **When** would it happen?

1. When events unevenly occur (e.g., rare cancer)
2. When data collection is uneven (e.g., only positive survey respondents)

# Issue: Imbalanced Data

## 4. How to check for it?

- Visualize **histogram** or **bar chart** of feature values



# Issue: Imbalanced Data

## 5. How to mitigate it?

- Collect more data instances
- Resample instances (e.g., Undersampling, Oversampling, SMOTE)

Further reading:


<https://machinelearningmastery.com/dimensionality-reduction-for-machine-learning/>

# Data Resampling


Image credit:

<https://www.analyticsvidhya.com/blog/2020/07/10-techniques-to-deal-with-class-imbalance-in-machine-learning/>

## Undersampling



## Oversampling



Data **leakage (snooping)**: remember to **first split** dataset to train–test, then **resample** train and test datasets separately

# Synthetic Minority Oversampling Technique (SMOTE)




Image Credit: <https://www.quora.com/Can-you-explain-me-SMOTE-Synthetic-Minority-Over-sampling-Technique-in-simple-terms>



*Questions!*







# Wrapping Up

# What did we learn this week?

## Data Issues

1. Linear Separability
2. Curse of Dimensionality
3. Imbalanced Data

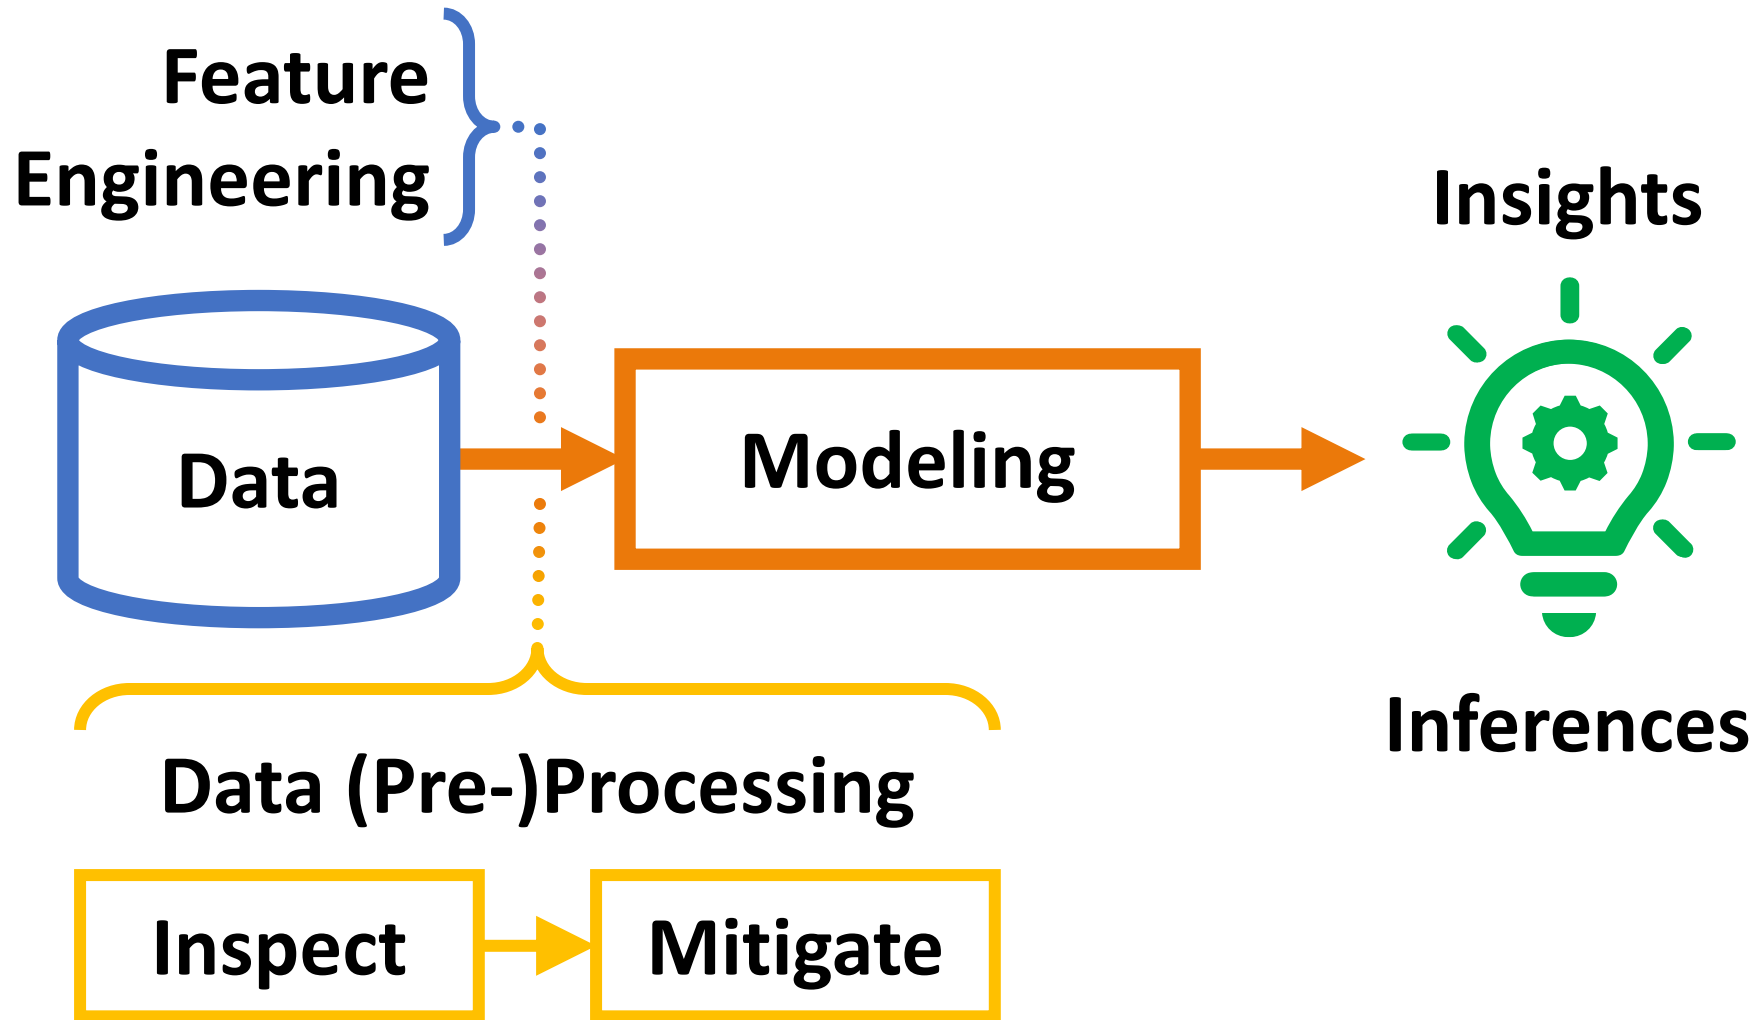
## Issue Template

1. **What** is the issue?
2. **Why** is it a problem?
3. **When** would it happen?
4. **How** to **check** for it?
5. **How** to **mitigate** it?

## Mitigations

1. Linear PCA, LDA  
(for Linear Separability, Dimensionality Reduction)
2. Feature Selection  
(Recursive Feature Elimination, Correlation, Mutual Information)
3. Resampling  
(Undersampling, Oversampling, SMOTE)

# Machine Learning Pipeline





# On Thursday: Feature Engineering