

W07-12 Recap + Final Exam Practice Review

13

B

CS 3244
Machine Learning




NUS | Computing

W07-12 Recap

Classification

$y \in \{0,1\}$ binary
 $y \in \{y_A, y_B, \dots\}$ multi-class



$y = M(x), \quad x = \vec{x} = (x_1, x_2)^T$

[W07b] Student Learning Outcomes

What did we learn for Evaluation?

1. Classification vs. Regression

2. Classification Metrics

3. Regression Metrics

1. Accuracy

2. Confusion Matrix, TP, TN, FP, FN

3. Precision, Recall, F_1

4. ROC, AUC

5. Micro- and Macro-Averaging

6. PR-AUC (Average Precision)

1. 1D regression: MSE, MAE

2. Vector regression: Euclidean distance, Angular distance / Cosine Similarity

Exercise E07.1

Match appropriate **evaluation metric** to challenge

Challenge	Evaluation Metric
Imbalanced actual classes	Accuracy (Emote :one:)
	Precision (:two:)
	Recall (:three:)
Multiclass classification	F_1 Score (:four:)
	ROC AUC (:five:)
	PRC AUC (:six:)
Cost-dependent classes	Micro-Average (:seven:)
	Macro-Average (:eight:)


Emote (react) in Slack #general channel one or more options (MRQ) for each challenge

NUS CS3244: Machine Learning

17

Regression

$y \in \mathbb{R}$ any real number



$y = M(x), \quad x = x_1$

Classification Metrics

		Actual Label		
		Alert	Not	
Predicted Label	Alert	<div>2</div> <div>True Positive</div>	<div>1</div> <div>False Positive</div>	<div>3</div> <div>Σ Pred. Pos.</div>
	Not	<div>3</div> <div>False Negative</div>	<div>4</div> <div>True Negative</div>	<div>7</div> <div>Σ Pred. Neg.</div>
		<div>5</div> <div>Σ Actual Pos.</div>	<div>5</div> <div>Σ Actual Neg.</div>	

Precision


$P = TP / (TP+FP)$

F_1 Score

$F_1 = \frac{2}{\frac{1}{P} + \frac{1}{R}}$

Recall

$R = TP / (TP+FN)$



AUC is a **concise metric** instead of a full figure.
Concise metrics enable *clearer comparisons*.
AUC > 0.5 means the model is better than chance.
AUC \approx 1 means model is very accurate.

Regression Metrics

Average difference metrics

Mean Absolute Error (MAE)

$$MAE = \frac{1}{m} \sum_{j=1}^m |\hat{y}_j - y_j|$$

Mean Squared Error (MSE)

$$MSE = \frac{1}{m} \sum_{j=1}^m (\hat{y}_j - y_j)^2$$

Root Mean Squared Error (RMSE)

$$RMSE = \sqrt{\frac{1}{m} \sum_{j=1}^m (\hat{y}_j - y_j)^2}$$

MSE and RMSE **penalize larger differences more** than MAE

Vector Distances and Similarity

Euclidean Distance

$$d = \sqrt{(\hat{y} - y)^T (\hat{y} - y)}$$

Dot Product

Cosine Similarity

$$s = \cos(\theta) = \frac{\hat{y} \cdot y}{\|\hat{y}\| \|y\|}$$

Angular Distance

$$\theta = \cos^{-1}(s)$$

Student Learning Outcomes

What did we learn this week?

Data Issues

- 1. [Linear Separability](#)
- 2. [Curse of Dimensionality](#)
- 3. [Imbalanced Data](#)

Issue Template

- 1. **What** is the issue?
- 2. **Why** is it a problem?
- 3. **When** would it happen?
- 4. **How** to check for it?
- 5. **How** to mitigate it?

Checks

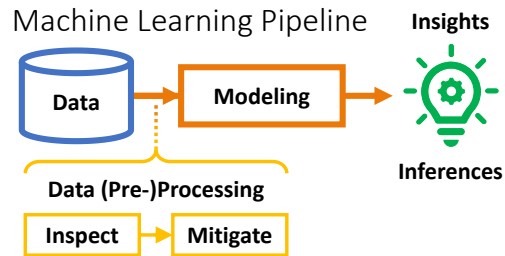
- 1. [Linear SVM](#), [PCA](#), [LDA](#)
- 2. [Visualize Histograms](#)

Mitigations

- 1. [Dimensionality Reduction](#) ([PCA](#), [LDA](#), [Deep Auto-Encoders](#))
- 2. [Feature Selection](#) ([Recursive Feature Elimination](#), [Correlation](#), [Mutual Information](#))
- 3. [Resampling](#) ([Under/Oversampling](#), [SMOTE](#))


NUS CS3244: Machine Learning

52




Issue: Linear Separability

Linearly Separable?



Not without data processing



How to make linearly separable?


$$x = \begin{pmatrix} x_1 \\ x_2 \end{pmatrix}$$

$$x' = \begin{pmatrix} (x_1 - \bar{x}_1)^2 \\ (x_2 - \bar{x}_2)^2 \end{pmatrix} = (x - \bar{x})^T (x - \bar{x})$$

Feature Engineering!

Testing Linear Separability with Linear Soft-Margin SVM

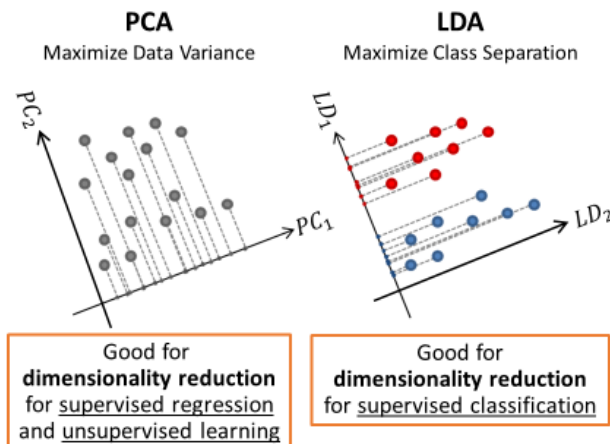
- Each $\xi^{(j)}$ is the distance that the misclassified point j is from its correct margin
- Total violation: $\sum_{j=1}^m \xi^{(j)}$
- Calculating the total violation indicates how linearly separable the data is in terms of its features
- Higher violation => Less linearly separable



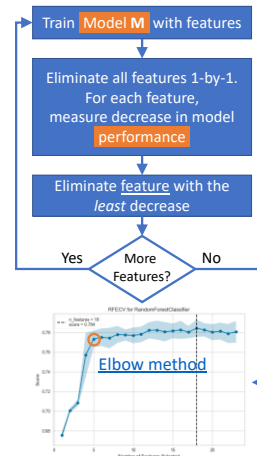
Issue: Curse of Dimensionality

Dimensionality Reduction

Feature Selection



Recursive Feature Elimination



Information gain

A chosen feature x_i divides the example set S into subsets S_1, S_2, \dots, S_c according to the C_i distinct values for x_i . The entropy then reduces to the entropy of the subsets S_1, S_2, \dots, S_c :

$$\text{remainder}(S, x_i) = \sum_{j=1}^{C_i} \frac{|S_j|}{|S|} H(S_j)$$


Information Gain (IG; "reduction in entropy") from knowing the value of x_i . Choose the attribute with the largest IG:

$$\text{IG}(S, x_i) = H(S) - \text{remainder}(S, x_i)$$

NUS CS3244: Machine Learning



Issue: Imbalanced Data

Oversampling

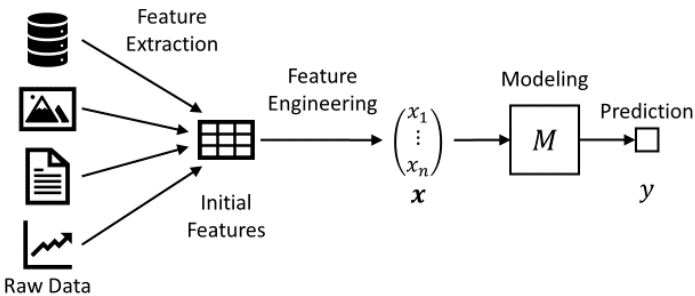


SMOTE

- Steps
- 1. Consider minority and majority instances in vector space.
- 2. For each minority-class instance pair, interpolate their feature values.
- 3. Randomly synthesize instances and label with minority class
- 4. More instances added to minority class



Feature Extraction/Engineering → Modeling

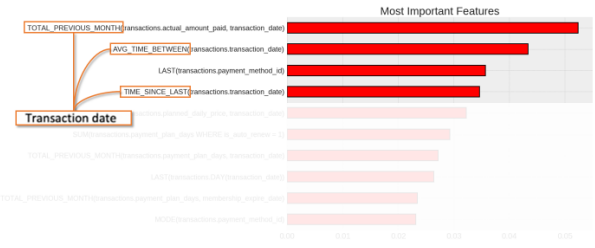


NUS CS3244: Machine Learning

9

Tabular

Tabular Feature Engineering:
Counting, Aggregation, Difference, Min, Max



Source: <https://github.com/featuretools/predict-customer-churn>

NUS CS3244: Machine Learning

4

Temporal

Sliding Time Window

- Prediction Task: **Price Prediction**
- Features
 - Moving Average
 - Moving Standard Deviation
 - Moving Range (Min, Max)
 - Moving Trend (Slope of linear fit)




Image credit: <https://cloud.google.com/blog/products/ai-machine-learning/how-to-quickly-solve-machine-learning-forecasting-problems-using-pandas-and-bigquery>

NUS CS3244: Machine Learning

19

[W08b] Student Learning Outcomes

What did we learn?

1. Describe **issues** when extracting features for various data types
2. Describe **techniques** of feature extraction/engineering for different data types

Tabular	Temporal	Image	Text
<ul style="list-style-type: none">• Domain-specific custom equations• Features from counting, aggregation, difference, min, max	<ul style="list-style-type: none">• Features from previous values, aggregate statistics, linear regression• Wave analysis features	<ul style="list-style-type: none">• RGB image as 3D tensor• Color features from RGB histogram• Shape features from edge detection• Edge detection via Convolution	<ul style="list-style-type: none">• Tokenization• Stemming, Lemmatization• Stop words• Bag-of-Words encoding

NUS CS3244: Machine Learning

52

Image

Feature: Edge Detection Kernels

$$\frac{\partial c}{\partial p_x} \approx \frac{c(x+1,y) - c(x-1,y)}{p(x+1,y) - p(x-1,y)}$$

$$I_{p_x} = \begin{pmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{pmatrix} \quad x = \begin{pmatrix} c(-1,-1) & c(0,-1) & c(1,-1) \\ c(-1,0) & c(0,0) & c(1,0) \\ c(-1,1) & c(0,1) & c(1,1) \end{pmatrix} \quad I_{p_x} * x = \begin{pmatrix} -c(-1,-1) + 0 + c(1,-1) \\ -c(-1,0) + 0 + c(1,0) \\ -c(-1,1) + 0 + c(1,1) \end{pmatrix}$$

$$I_{p_x} \text{ is a Convolution Matrix (Kernel)}$$

$$x = \begin{pmatrix} 9 & 9 & 3 & 3 & 4 \\ 9 & 3 & 3 & 4 & 5 \\ 9 & 3 & 3 & 4 & 5 \\ 9 & 3 & 3 & 4 & 5 \\ 9 & 3 & 3 & 4 & 5 \end{pmatrix} \quad I_{p_x} * x = \begin{pmatrix} -6 - 6 - 6 & -6 + 1 + 2 & 1 + 2 + 2 \\ -6 - 6 - 6 & 1 + 2 + 1 & 2 + 2 + 2 \\ -6 - 6 - 6 & 2 + 1 - 6 & 2 + 2 + 1 \end{pmatrix}$$

$$* \text{ is Convolution operator}$$

$$\text{Means: element-wise multiply, then sum}$$

NUS CS3244: Machine Learning

32

Text

Bag-of-Words (BOW) Encoding

1. Preprocess string s to array of words w
2. Array of words \rightarrow One-hot vector (fixed length)
3. BOW(w) $\rightarrow x$
4. Problem: **high dimensions** if many words

$$x^{(1)} = \begin{pmatrix} 1 \\ 1 \\ 1 \\ 1 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \end{pmatrix} \quad x^{(2)} = \begin{pmatrix} 0 \\ 1 \\ 1 \\ 1 \\ 1 \\ 1 \\ 1 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \end{pmatrix}$$


#	Original Text s	Pre-Processed Words w	chicken	wings	amazing	honestly	but	way	too	long	wait	not	worth	salty	expensive
1	"Chicken wings were amazing honestly"	['chicken', 'wings', 'amazing', 'honestly']	1	1	1	1	0	0	0	0	0	0	0	0	0
2	"Amazing wings, but waaaay too long to wait."	['amazing', 'wings', 'but', 'way', 'too', 'long', 'wait']	0	1	1	0	1	1	1	1	0	0	0	0	0
3	"Not worth it! Too salty chicken and expensive!"	['not', 'worth', 'too', 'salty', 'chicken', 'expensive']	1	0	0	0	0	0	1	0	0	1	1	1	1

The word "too" could predict **negative** sentiment

46


Perceptron

Perceptron



NUS CS3244: Machine Learning

Perceptron Learning Algorithm

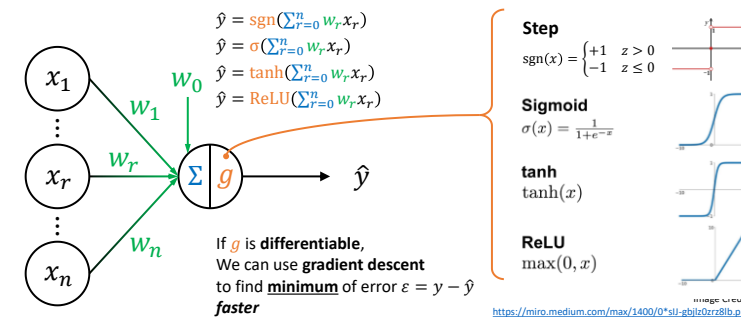
1. Initialize weights \mathbf{w}
 - Could be all zero, or random small values
 2. For each instance i with features $\mathbf{x}^{(i)}$
 - Classify $\hat{y}^{(i)} = \text{sgn}(\mathbf{w}^T \mathbf{x}^{(i)})$
 3. Select one **misclassified** instance
 - Update weights: $\mathbf{w} \leftarrow \mathbf{w} + \eta(\mathbf{y} - \hat{\mathbf{y}})\mathbf{x}$
 4. Iterate steps 2 to 3 until
 - Convergence (classification error < threshold), or
 - Maximum number of iterations
- 

$$\begin{pmatrix} w_0 \\ w_1 \\ \vdots \\ w_r \\ \vdots \\ w_n \end{pmatrix} \leftarrow \begin{pmatrix} w_0 \\ w_1 \\ \vdots \\ w_r \\ \vdots \\ w_n \end{pmatrix} + \eta(y - \hat{y}) \begin{pmatrix} 1 \\ x_1 \\ \vdots \\ x_r \\ \vdots \\ x_n \end{pmatrix}$$

$$w_r \leftarrow w_r + \eta(y - \hat{y})x_r$$

NUS CS3244: Machine Learning

Differentiable Activation Functions




https://miro.medium.com/max/1400/0*sIJ-gbjlz0zrz8lb.p

NUS CS3244: Machine Learning

30

Neural Networks

Multi-Layer Perceptron (Neural Network)



NUS CS3244: Machine Learning

Chain Rule

Consider composite function

$$g(x) = g(f(x))$$

$$g = g(f), f = f(x)$$

$$g'(x) = \frac{dg}{dx} = \frac{dg}{d\textcolor{blue}{f}} \frac{d\textcolor{blue}{f}}{dx}$$

Intuition

Rate of change of g relative to x is the product of

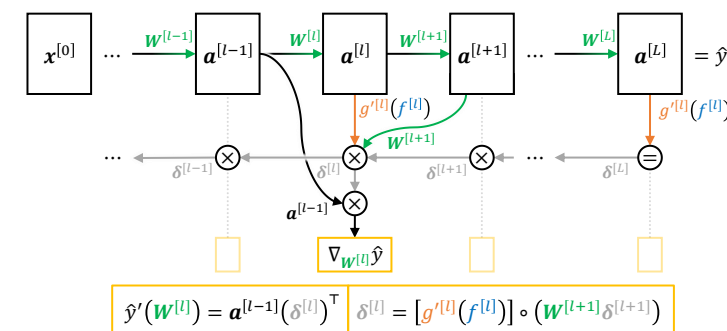
- rates of change of g relative to f and
- rates of change of f relative to x

"If

- a car travels 2x fast as a bicycle and
 - the bicycle is 4x as fast as a walking man,
- then the car travels $2 \times 4 = 8$ times as fast as the man.”
- George F. Simmons, *Calculus with Analytic Geometry* (1985)

NUS CS3244: Machine Learning

Backward Propagation



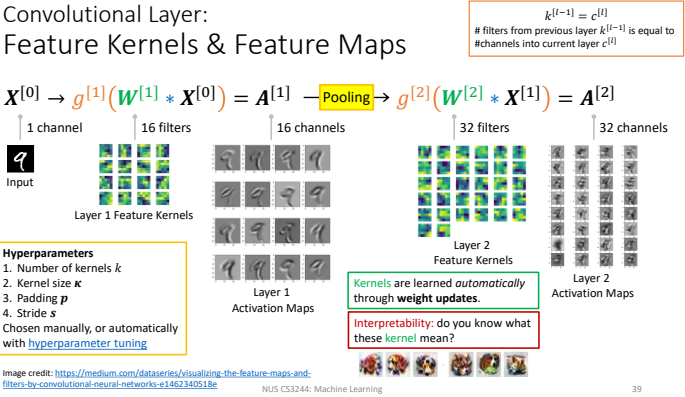

NUS CS3244: Machine Learning

23

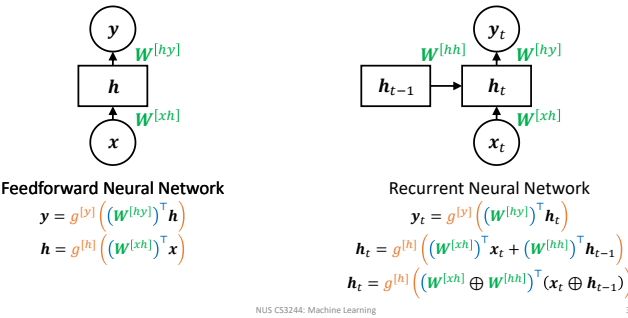
CNN

RNN

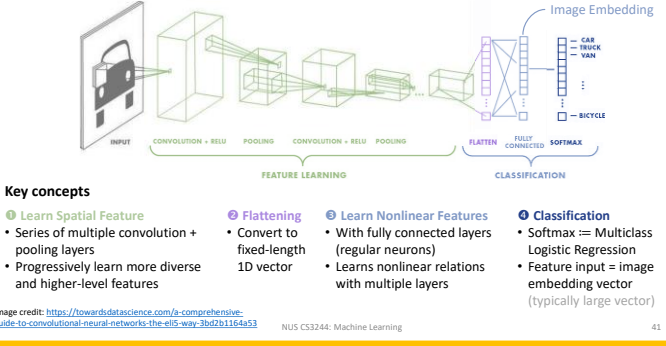
From Manual Feature Engineering To Architecture Engineering



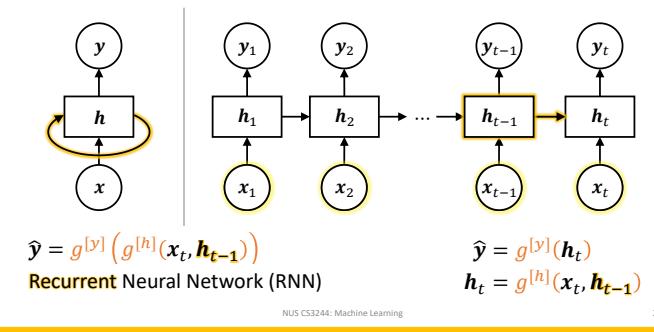
RNN Weights



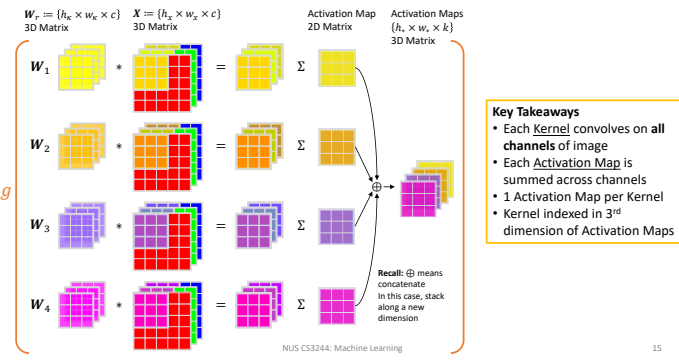
Convolutional Neural Network



Neurons with Recurrence



Multi-Channel Convolutions ($c = 3$ channels, $k = 4$ filters)



Training Issues


Deep Learning Training Issues → Mitigations

- Overfitting → Dropouts
- Saturating Gradient → ReLU Activation Function
- Vanishing Gradients → ResNet “Shortcuts”, LSTM “Forget” Gates

[W09a] Student Learning Outcomes

What did we learn?

- Feature Engineering → Architecture Engineering
- CNN: exploits spatial information using convolutions
- RNN: exploits history information using recurrence



Interpretable
“Glassbox” Models

Exercise E11a.1 Solution

How would you interpret?
Linear Regression

$$\hat{y} = w_0 + w_1x_1 + w_2x_2 + \dots + w_nx_n = \sum_{r=0}^n w_r x_r$$
$$= \mathbf{w} \cdot \mathbf{x} = \mathbf{w}^T \mathbf{x}$$

$\hat{y} \propto w_r x_r, \forall r$ If $w_1 = k w_2$, then x_1 is k times more important than x_2

Weighted Sum Interpretation

- Bigger w_r means
- Larger weight
 - More importance for x_r
 - Direction? Supportive (positive) or opposing (negative) influence

Gradient Interpretation

- Bigger w_r means
- Steeper slope for x_r axis
 - Changes in x_r lead to bigger in \hat{y} changes
 - More importance for x_r
 - Direction indicates increasing or decreasing influence

NUS CS3244: Machine Learning

12

Exercise E11a.2 Solution

How would you interpret?
Logistic Regression

$$\hat{y} = \sigma(z) = \frac{1}{1 + e^{-z}}$$

$$f = \mathbf{w} \cdot \mathbf{x} = \sum_{r=0}^n w_r x_r$$

$\text{logit}(P(\hat{y})) \propto w_r x_r, \forall r$ If $w_1 = k w_2$, then log odds ratio of x_1 is k times bigger than of x_2

Weighted Sum Interpretation

- Bigger w_r means
- Larger importance
 - Direction indicates influence

Gradient Interpretation


- Bigger w_r means?
- Steeper slope for x_r near decision boundary
 - Decision boundary more perpendicular to x_r
 - Weight sign indicates direction of pos/neg prediction

NUS CS3244: Machine Learning

23

LIME
Local, Model-Agnostic Explanations

LIME
Local Interpretable Model-agnostic Explanations



Prediction Model

- $f(\mathbf{x})$
- Non-linear model of $\mathbf{x} = \begin{pmatrix} x_1 \\ x_2 \end{pmatrix}$
 - Shown as curvy decision boundary

Explanation: LIME

1. Starting with instance \mathbf{x} to explain
2. Focus on Local region
3. Training set as neighbors $\mathbf{x}^{(\eta)} \in \mathcal{X}^{(\eta)}$
4. Train surrogate model, e.g., linear:
$$g(\mathbf{x}) = \mathbf{w} \cdot \mathbf{x} = \sum_{r=0}^n w_r x_r$$

Image Credit: <https://santiagof.medium.com/model-interpretability-making-your-model-confess-lime-89db7f70a72b>

NUS CS3244: Machine Learning

28

LIME
Find “best” explainer g that minimizes $\xi(\mathbf{x})$

Python API:

<https://github.com/marcotcr/lime>

$$\xi(\mathbf{x}) = \underset{g \in G}{\text{argmin}} \left(\overset{\text{“Faithful”}}{L(f, g, \pi_x)} + \overset{\text{“Simple”}}{\Omega(g)} \right)$$

- f is the predictor function (model)
 g is the explainer function (model)
 $\pi_x(\mathbf{x}^{(\eta)})$ is the neighbor proximity function
- E.g., exponential decay $\exp(-d(\mathbf{x}, \mathbf{x}^{(\eta)}))^2$

$L(f, g, \pi_x)$ is the locally-weighted error loss function between predictor f and explainer g

$$L(f, g, \pi_x) = \sum_{\mathbf{x}^{(\eta)} \in \mathcal{X}^{(\eta)}} \pi_x(\mathbf{x}^{(\eta)}) (f(\mathbf{x}^{(\eta)}) - g(\mathbf{x}^{(\eta)}))^2$$

$\Omega(g)$ is the sparsity regularization

- Want simpler explanation
 - \Rightarrow fewer weights
 - \Rightarrow Lasso (L1 norm)
- Penalizes if total weights is too large
- λ is hyperparameter on how much to penalize

$$\Omega(g) = \lambda \|\mathbf{w}_r\|_1 = \lambda \sum_{r=1}^n w_r$$

Ribeiro, M. T., Singh, S., & Guestrin, C. (2016, August). “Why should I trust you?” Explaining the predictions of any classifier. KDD’16.

NUS CS3244: Machine Learning

29

Grad-CAM
Gradient-Weighted
Class Activation Maps

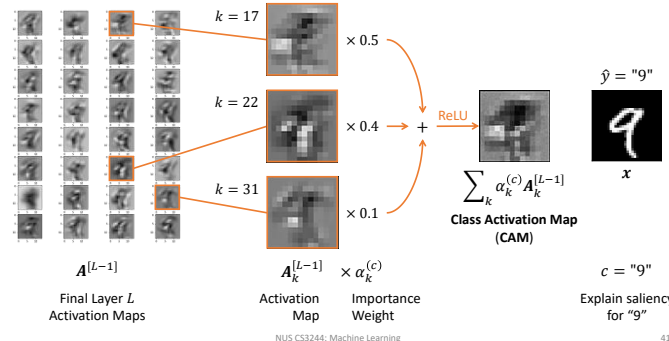
Grad-CAM Steps

1. Compute Activation Maps $\mathbf{A}^{[L]}$ of last conv layer L
 1. via Forward Propagation
2. Choose class label c to explain about (e.g., predict “9”, “car”)
3. Filter prediction \hat{y} to be about class c
 1. Given: $\hat{y} = \begin{pmatrix} \hat{y}^{(1)} \\ \hat{y}^{(2)} \\ \vdots \\ \hat{y}^{(n)} \end{pmatrix}$, $\mathbf{e}^{(c)} = \begin{pmatrix} 0 \\ 0 \\ \vdots \\ 1 \\ 0 \end{pmatrix}$, then $\tilde{\mathbf{y}}^{(c)} = \hat{\mathbf{y}} \circ \mathbf{e}^{(c)} = \begin{pmatrix} 0 \\ 0 \\ \vdots \\ y^c \\ 0 \end{pmatrix}$
 2. To generate explanation only for that class c
4. Compute importance weight $\alpha_k^{(c)}$ for each Activation Map $\mathbf{A}_k^{[L]}$
 1. Backprop from $\tilde{\mathbf{y}}^{(c)}$ to get gradients (relative to activations) at last conv layer
5. Compute weighted sum with ReLU to get Class Activation Map

NUS CS3244: Machine Learning

40

Grad-CAM example: Why did the CNN predict “9”?



NUS CS3244: Machine Learning


41

k-Means Clustering

k-Means Objective
Minimize **Within-Cluster Sum-of-Squares (WCSS)** (i.e. variance)

$$L = \arg \min_S \sum_{c=1}^k \sum_{x \in S_c} \|x - \mu_c\|^2$$

- $S = \{S_1, S_2, \dots, S_c, \dots, S_k\}$ is the set of all clusters
- k is the total number of clusters
- S_c is the c th cluster of points
- Note that c refers to cluster, not class
- $x \in S_c$ refers to a point in cluster S_c
- $\mu_c = \frac{1}{|S_c|} \sum_{x \in S_c} x$ is the centroid point in cluster S_c
- $\|x - \mu_c\|^2$ refers to the squared Euclidean distance from x to μ_c



NUS CS3244: Machine Learning 13

k-Means clustering algorithm

```
for c = 1 to k:
    μ_c ← Random() } 1) Initialize cluster centroids

while not Converged():
    for j = 1 to m:
        y^(j) ← c = argmin_c ||x^(j) - μ_c||^2 } 2) Assign datapoints to clusters
        x ← S_c
    for c = 1 to k:
        μ_c ← (1/|S_c|) Σ_{x ∈ S_c} x } 3) Update cluster centroids
    t += 1

return y } Cluster labels of all datapoints
```

NUS CS3244: Machine Learning 14

How to choose k (number of clusters)?

- Use domain knowledge
- Note the k with *diminishing return*
 - When k is too high, marginal decrease in **within-cluster sum-of-squares (WCSS)**
 - How to see?
 - “Elbow” method





Image credit: <https://www.analyticsvidhya.com/blog/2021/05/k-mean-getting-the-optimal-number-of-clusters/>

NUS CS3244: Machine Learning 17

Auto-Encoders

Auto-Encoder (with Bottleneck)




Reconstruction Loss Function
Squared Euclidean Distance
 $\epsilon_x(\hat{x}, x) = (\hat{x} - x)^T (\hat{x} - x)$

Benefits

- Compressed representation of x in the middle layer

NUS CS3244: Machine Learning 46

Auto-Encoder Training



Reconstruction Loss Function
Squared Euclidean Distance
 $\epsilon_x(\hat{x}, x) = (\hat{x} - x)^T (\hat{x} - x)$

Gradient

$$\frac{\partial \epsilon_x}{\partial W} = \frac{\partial \hat{x}}{\partial W} \frac{\partial \epsilon_x}{\partial \hat{x}} \quad \hat{y} := \hat{x}$$
$$\frac{\partial \hat{x}}{\partial W} = \frac{\partial D}{\partial W} = \frac{\partial E}{\partial W} \frac{\partial D}{\partial E} = \frac{\partial E}{\partial W} \frac{\partial D}{\partial \hat{z}}$$


$$\frac{\partial D}{\partial W^{[l_D]}} = \frac{\partial f^{[l_D]}}{\partial W^{[l_D]}} \dots \frac{\partial f^{[l_D]}}{\partial g^{[l_D-1]}} \frac{\partial g^{[l_D]}}{\partial f^{[l_D]}}$$
$$\frac{\partial E}{\partial W^{[l_E]}} = \frac{\partial f^{[l_E]}}{\partial W^{[l_E]}} \dots \frac{\partial f^{[l_E]}}{\partial g^{[l_E-1]}} \frac{\partial g^{[l_E]}}{\partial f^{[l_E]}}$$
$$\frac{\partial D}{\partial W^{[l_E]}} = \frac{\partial E}{\partial W^{[l_E]}} \frac{\partial D}{\partial \hat{z}}$$

Take-Away:
Backprop gradient descent for weight update
 $W \leftarrow W - \eta \frac{\partial \epsilon_x}{\partial W}$
Same as all neural networks, but through 2 models

$$l_E < l_D$$

NUS CS3244: Machine Learning 55

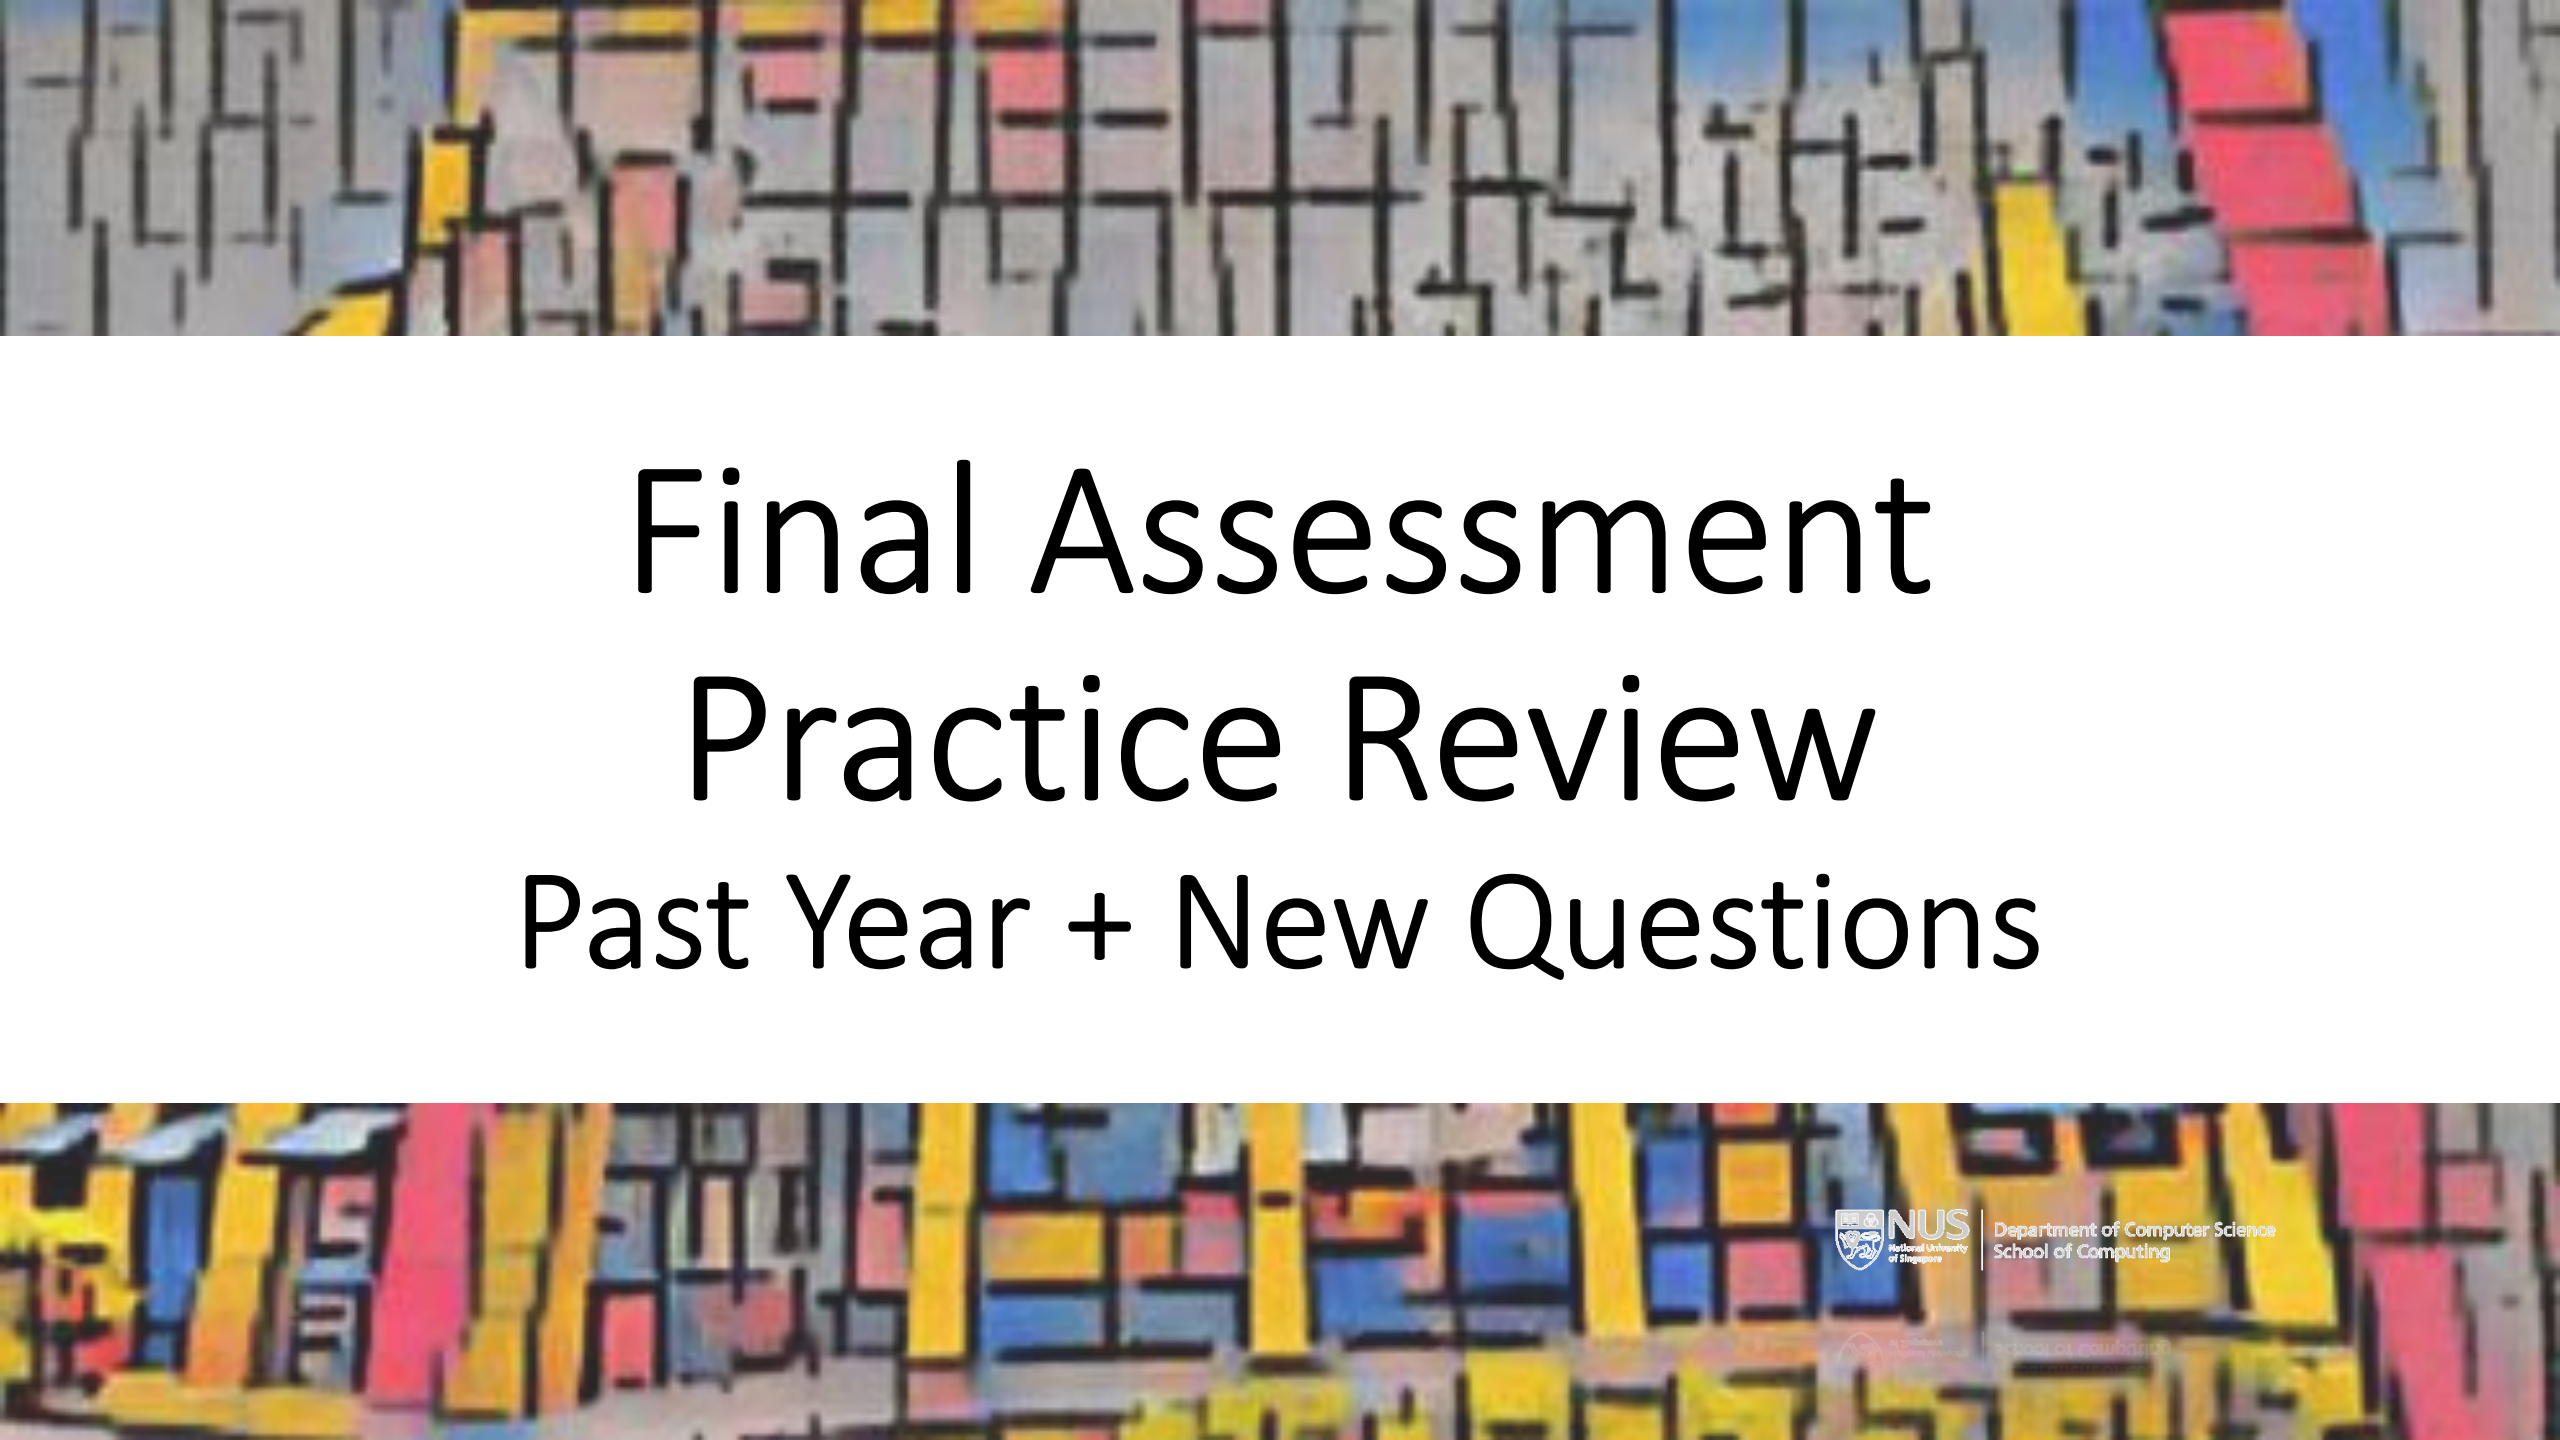
Auto-Encoder for Feature Representation Learning



Reconstruction Loss Function
Squared Euclidean Distance
 $\epsilon_x(\hat{x}, x) = (\hat{x} - x)^T (\hat{x} - x)$

Prediction Loss Function
Cross-Entropy (Classification)
 $\epsilon_y(\hat{y}, y) = -y^T \log \hat{y}$
Squared Error (Regression)
 $\epsilon_y(\hat{y}, y) = (\hat{y} - y)^T (\hat{y} - y)$

t-SNE projection of latent space \hat{z} representations



Final Assessment Practice Review

Past Year + New Questions



Department of Computer Science
School of Computing



School of Computing




Bonus (Optional) AMA

Reading Week, 17 Nov, Wed @ 4-6pm



Department of Computer Science
School of Computing



School of Computing

Goodbye and Good Luck

Final Exam

24 Nov (Wed) 16:00-18:00.

Venue: Zoom Proctor Rooms and also
Seminar Room 1 (COM1 #02-06; SR1)

Physical, online notes, (online) calculator.
Not open internet / Web.

Stay in touch with our [CS3244 alumni](#)
and general interest group on
Facebook (search for “cs3244”)



Thanks for joining us for the journey!

Photo Credits: [People's Online Daily](#) (English Version)